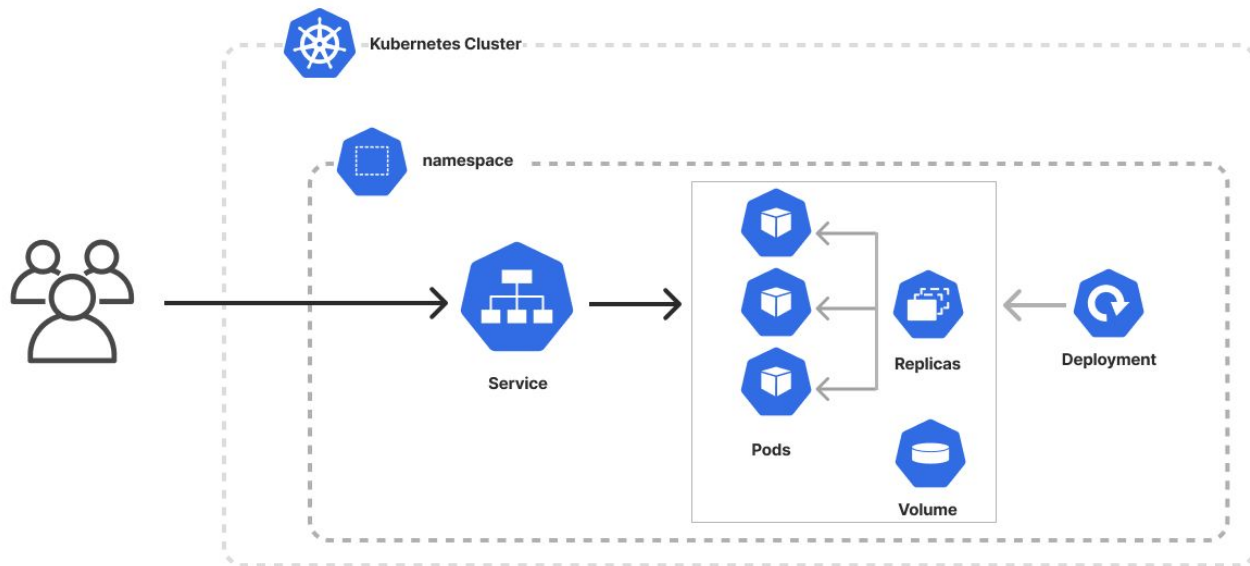




Troubleshooting Mastery: DevOps Solutions for Java OOM Challenges

Daksh Jain

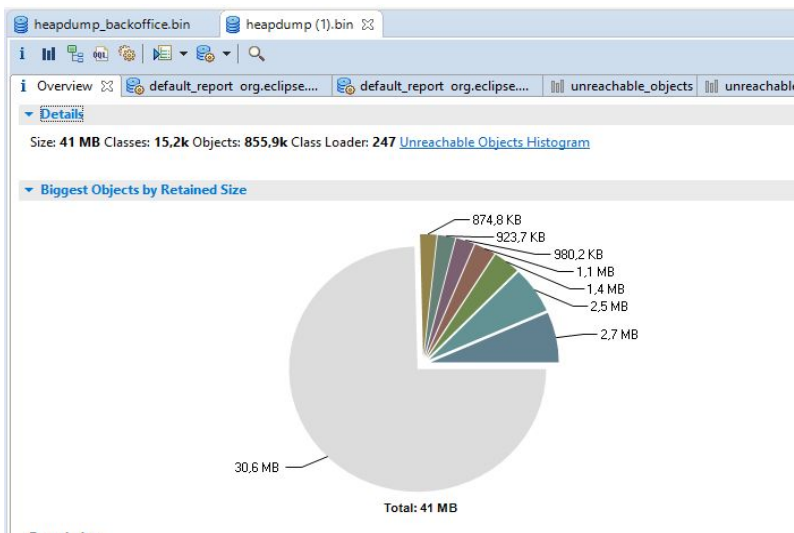
Setting the context...



```

run:
Performing 10000000 append operations;process completed in :129ms
Performing 20000000 append operations;process completed in :271ms
Performing 30000000 append operations;process completed in :495ms
Performing 40000000 append operations;process completed in :509ms
Performing 50000000 append operations;process completed in :860ms
Performing 60000000 append operations;process completed in :950ms
Performing 70000000 append operations;process completed in :1025ms
Performing 80000000 append operations;process completed in :1051ms
Performing 90000000 append operations;process completed in :1071ms
Exception in thread "main" java.lang.OutOfMemoryError: Java heap space
    at java.util.Arrays.copyOf(Arrays.java:3332)
    at java.lang.AbstractStringBuilder.ensureCapacityInternal(AbstractStr
    at java.lang.AbstractStringBuilder.append(AbstractStringBuilder.java:
    at java.lang.StringBuilder.append(StringBuilder.java:136)
    at com.day08.stringmanipulation.SpeedDemoClass.iterator(SpeedDemoClas
    at com.day08.stringmanipulation.SpeedDemoClass.main(SpeedDemoClass.ja

```




Problems

#1 Cost

- Enterprise tools offer robust monitoring and debugging features but come with **high costs**.
- Startups and larger companies prioritize **cost optimization**, making these solutions less attractive:
 - **Budget Constraints:** Startups and smaller companies often have **limited budgets** and prefer to allocate resources efficiently.
 - **In-House Expertise:** Companies with strong tech teams want to develop **custom solutions tailored to their specific needs**, reducing reliance on expensive third-party tools.



Problems



#2 Dependency

Even if Open-source solutions are setup, it often require DevOps oncall to provide heap dump files to developers **manually**.

This process is operation-intensive, leading to delays and inefficiencies.

Our actions should be process dependant not person dependant 🚀





Before jumping in...


After today's discussion what I want is, if a Dev team receives a pager for production apps and it has OOM error then -

- Critical files like heap dumps and thread dumps should be **auto created**.
- These files must be **available** at the time of error to **reduce MTTR**.
- Devs have **access** to these files while debugging.



Final actionables

Collect & provide heap dump from apps running in K8S pods.



Find a solution to automate heap dump creation, collection & delivery to respective team 24x7.



FAQs

What is heap dump ?


A heap dump is a snapshot of all the objects in the Java Virtual Machine (JVM) heap at a certain point in time. The JVM software allocates memory for objects from the heap for all class instances and arrays.

What is thread dump ?

A thread dump is a snapshot of the state of all the threads of a Java process. The state of each thread is presented with a stack trace, showing the content of a thread's stack.

OOM Status code in K8S ?

The Kubernetes OOMKilled (Exit Code **137**) is a signal sent by the Linux Kernel to terminate a process due to an Out Of Memory (OOM) condition. This event is usually an indication that a container in a pod has exceeded its memory limit and the system cannot allocate additional memory.



**So what are the possible ways we can help
Devs perform seamless debugging?**



Solution #1

Use [Lifecycle hook](#) with pre stop setting in the deployment manifest.

```
lifecycleHooks:  
  heapDumpCollectionEnabled: true  
  settings:  
    preStop:  
      exec:  
        command: ["/bin/bash", "heapDump.sh", "unique-s3-bucket-name", "s3-bucket-region"]
```



```

apiVersion: v1
kind: ConfigMap
metadata:
  name: {{ template "backend.fullname" . }}-heap-configmap
  labels:
    app: {{ template "backend.name" . }}
    chart: {{ template "backend.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
data:
  heapDump.sh: |-
    #!/bin/bash
    set -e
    BUCKET_NAME=$1
    BUCKET_REGION=$2
    NAMESPACE={{ .Release.Namespace }}
    AWS_BIN='which aws`
    SRC_DIR='/heapdump/log/'
    SRC_FILE="heapDump-`hostname`-`date +%F-%H-%M-%S`.hdprof"
    DST_FILE="s3://$BUCKET_NAME/$NAMESPACE/`hostname|awk -F- '{print $1}'`-`date +%F`/`hostname`-`date +%F-%H-%M-%S`.hdprof"
    PID=1
    unset JAVA_TOOL_OPTIONS
    export JAVA_TOOL_OPTIONS='-Xmx128m'
    mkdir -p "$SRC_DIR"

    if [[ -f "$SRC_FILE" ]]; then
      echo 'Begin heap dump upload'
      $AWS_BIN s3 cp "$SRC_DIR/$SRC_FILE" "$DST_FILE" --region "$BUCKET_REGION"
      rm -rf "$SRC_DIR/$SRC_FILE"
    else
      jmap -dump:live,format=b,file="$SRC_DIR/$SRC_FILE" ${PID}
      $AWS_BIN s3 cp "$SRC_DIR/$SRC_FILE" "$DST_FILE" --region "$BUCKET_REGION"
      rm -rf "$SRC_DIR/$SRC_FILE"
    fi


```

- Input
 - Bucket name
 - Region


- Steps
 - Run **jmap** cmd to collect heap dump
 - Run **aws s3 cp** to copy file to S3 bucket.

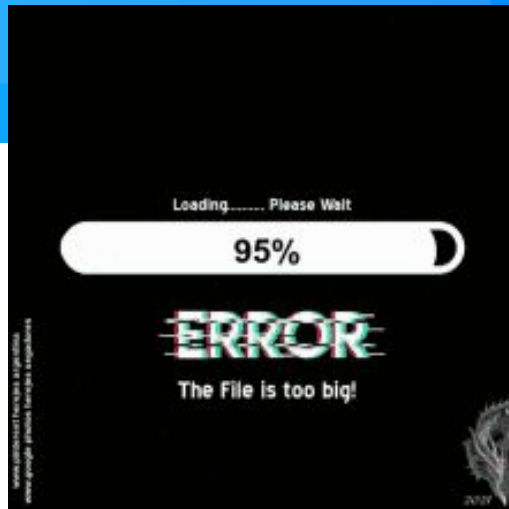


Issue in Solution #1

- “Helpful” heap dumps are usually **GBs** in size. 
- To transfer it to a S3 bucket we run **aws s3 cp**.
- Now the container will be killed if -

`terminationGracePeriodSeconds < hook run time + container stop time`

- In this case the hook run time is a **variable** + it is a heavy command to execute.
- So we are not able to determine the **exact value** for grace period.
- Eventually we will **lose** the heap dump. 



Solution #2



Run the same heap dump script in **JAVA_TOOL_OPTIONS**.

This helps us to remove the dependency on K8S and is **sort of a shift left approach** to the problem itself as we make the heap dump collection closer to the application.



```
JAVA_TOOL_OPTIONS: '-XX:+HeapDumpOnOutOfMemoryError  
-XX:HeapDumpPath=/data/backendHeapDump.hdprof  
-XX:OnOutOfMemoryError="mv /data/backendHeapDump.hdprof  
/data/`hostname`-`TZ=Asia/Kolkata date +%F-%H-%M-%S`.hdprof ;  
/usr/local/bin/aws s3 cp /data/`hostname`-`TZ=Asia/Kolkata date +%F-%H-%M-%S`.hdprof  
unique-bucket-url/path/`TZ=Asia/Kolkata date +%F`/`hostname`-`TZ=Asia/Kolkata date  
+%F-%H-%M-%S`.hdprof --region s3-bucket-region"'
```



Issue in Solution #2

But again same issue that if the container is killed then we lose the heap dump.



Solution #3

Use **persistent storage** instead of ephemeral storage for heap dump.



Implementation Pointers

✗ Don't use EBS.

Use EFS, to have support of **multi mount**.

In EFS use a **lifecycle rule** to move data to S3.

This removes all kinds of dependency –

- ✓ App **collects** heap dump on OOM
- ✓ Heap dump is **persistent** in EFS
- ✓ Move to S3 to reduce EFS cost + increased **accessibility**
- ✓ Devs can debug **independently** + **reduce MTTR**



aws

Services

Search

[Option+S]

Elastic Kubernetes Service

S3

Security Hub

IAM Identity Center

AWS Config

VPC

EC2

IAM

Elastic File System

File systems

Access points

AWS Backup

General

Performance mode

General Purpose

Throughput mode

Bursting

Lifecycle management

Transition into Infrequent Access (IA): 7 day(s) since last access

Transition into Archive: None

Transition into Standard: None

Availability zone

Regional

Amazon Elastic File System

Scalable, elastic, cloud-native NFS file system

Amazon Elastic File System (Amazon EFS) provides a simple, scalable, elastic file system for general purpose workloads for use with AWS Cloud services and on-premises resources.

Create file system

Create an EFS file system with recommended settings.

Create file system

Pricing

Access points (1)

Search access points by name or ID

Name

Access point ID

Path

POSIX user

Creation info

State

-

fsap-072b61cd6541cb32f

/srv/heapdump


1001 : 999

1001 : 999 (755)

Available

19

© 2024 Cloud Native Computing Foundation



```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: test-sc
provisioner: efs.csi.aws.com
reclaimPolicy: Delete
volumeBindingMode: Immediate
```





```
apiVersion: v1
kind: PersistentVolume

metadata:
  name: test-sc-pv
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  storageClassName: test-sc
  csi:
    driver: efs.csi.aws.com
    volumeHandle: fs-AAA::fsap-BBB
```



```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  namespace: production
  name: test-sc-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: test-sc
  resources:
    requests:
      storage: 5Gi
```



```

apiVersion: v1
kind: Pod
metadata:
  name: efs-setup
spec:
  securityContext:
    fsGroup: 999
    runAsGroup: 999
    runAsUser: 1001
    runAsNonRoot: true
  containers:
  - volumeMounts:
    - name: efs-storage
      mountPath: /srv/heapdump
    resources:
      requests:
        cpu: 0.1
        memory: 256Mi
      limits:
        cpu: 0.1
        memory: 256Mi
    image: <your image name here>
    imagePullPolicy: IfNotPresent
    name: disk-checked
    command: ["/bin/sh"]
    args: ["-c", "sleep 10000"]
  volumes:
  - name: efs-storage
    persistentVolumeClaim:
      claimName: efs-heapdump-test


```



<div> <input type="text"/> <div> <div>↻</div> <div>View details</div> <div>Delete</div> </div> </div>			
Path	POSIX user	Creation info	State
/srv/heapdump	1001 : 999	1001 : 999 (755)	🟢 Available



Learnings from Solution #3



We will use a path in our pod to mount the EFS - **/srv/data**. This path is first created as an access point in EFS.

Amazon EFS access points are **application specific entry points** into an EFS file system that make it easier to manage application access to shared datasets.

Once it is created by a particular user in AWS, this is a service level limitation that you **cannot** mount the **same path** by a **different user** in the **same EFS**.



