

ElasticSearch

Under The Hood

What is Big Data ?

And how do you manage it ?

What are Distributed Systems ?

What are MPP systems ?

Monolithic Systems

- Shared Memory Architecture

Distributed Systems

- Shared Disk Architecture
- Shared Nothing Architecture

Inverted Index

Document 1

The bright blue butterfly hangs on the breeze.

Document 2

It's best to forget the great sky and to retire from every wind.

Document 3

Under blue sky, in bright sunlight, one need not search around.

Document ID

Document

Stopword list

a
and
around
every
for
from
in
is
it
not
on
one
the
to
under

Inverted index

ID	Term	Document
1	best	2
2	blue	1, 3
3	bright	1, 3
4	butterfly	1
5	breeze	1
6	forget	2
7	great	2
8	hangs	1
9	need	3
10	retire	2
11	search	3
12	sky	2, 3
13	wind	2

Term
Ordinal

Terms
Dictionary

Postings
List

What is Lucene ?

Horizontally Scalable

Big Data

Shared Nothing Architecture

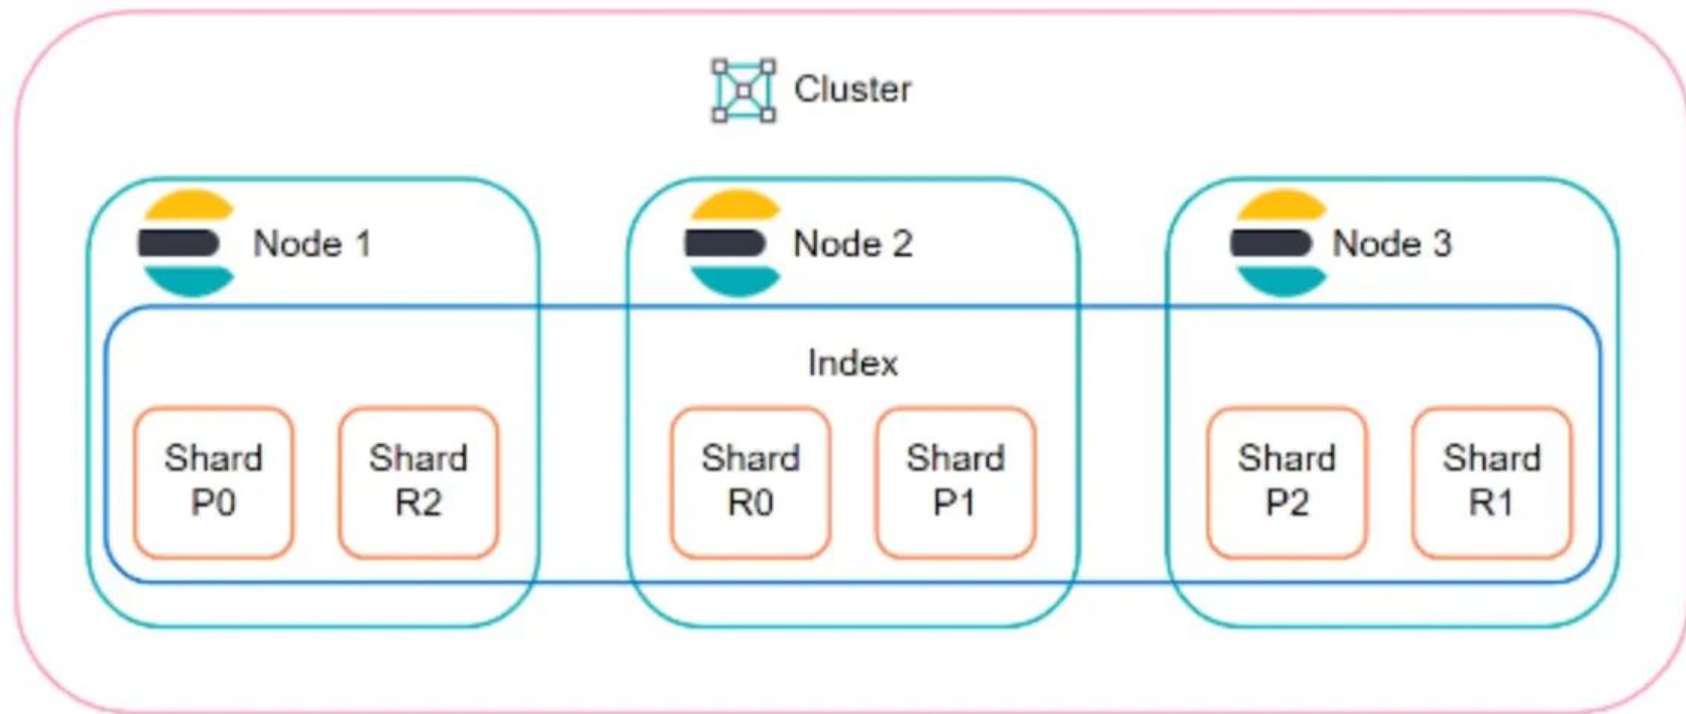
What is ElasticSearch ?

Search Engine

MPP

Distributed Database

Basic Data Spread Overview



Elasticsearch Index

Elasticsearch
shard

Elasticsearch
shard

Elasticsearch
shard

Elasticsearch
shard

Lucene index

Lucene index

Lucene index

Lucene index

Segment

Segment

Segment

Segment

Segment

Segment

Segment

Segment

We're starting the bottom up

We will come back to Elasticsearch


The Challenge Of Reliable Storage

Disks (both HDD/SSD)

The Challenge Of Reliable Storage

- Disks have **Disk Controllers**





The Challenge Of Reliable Storage

- Disks have **Disk Controllers** and they have **Caches**

The Challenge Of Reliable Storage

- Disks have **Disk Controllers** and they have **Caches INSIDE OF THEM**

The Challenge Of Reliable Storage

- Disks have **Disk Controllers** and they have **Caches INSIDE OF THEM**
- Disk IOPS Are Slow

Latency Comparison Numbers (~2012)

L1 cache reference	0.5	ns			
Branch mispredict	5	ns			
L2 cache reference	7	ns			14x L1 cache
Mutex lock/unlock	25	ns			
Main memory reference	100	ns			20x L2 cache, 200x L1 cache
Compress 1K bytes with Zippy	3,000	ns	3	us	
Send 1K bytes over 1 Gbps network	10,000	ns	10	us	
Read 4K randomly from SSD*	150,000	ns	150	us	~1GB/sec SSD
Read 1 MB sequentially from memory	250,000	ns	250	us	
Round trip within same datacenter	500,000	ns	500	us	
Read 1 MB sequentially from SSD*	1,000,000	ns	1,000	us	1 ms ~1GB/sec SSD, 4X memory
Disk seek	10,000,000	ns	10,000	us	10 ms 20x datacenter roundtrip
Read 1 MB sequentially from disk	20,000,000	ns	20,000	us	20 ms 80x memory, 20X SSD
Send packet CA->Netherlands->CA	150,000,000	ns	150,000	us	150 ms

Notes

1 ns = 10^{-9} seconds

1 us = 10^{-6} seconds = 1,000 ns

1 ms = 10^{-3} seconds = 1,000 us = 1,000,000 ns

Credit

By Jeff Dean: <http://research.google.com/people/jeff/>

Originally by Peter Norvig: <http://norvig.com/21-days.html#answers>

The Challenge Of Reliable Storage

- Have **Disk Controllers** and they have **Caches INSIDE OF THEM**
- Disk IOPS Are Slow
- Most databases rely on filesystems and not use disks directly

The Challenge Of Reliable Storage

- Have **Disk Controllers** and they have **Caches INSIDE OF THEM**
- Disk IOPS Are Slow
- Most databases rely on filesystems and not use disks directly

The Challenge Of Reliable Storage

- Have **Disk Controllers** and they have **Caches INSIDE OF THEM**
- Disk IOPS Are Slow
- Most databases rely on filesystems and not use disks directly
- And filesystems have their own metadata to maintain


```
#touch file.txt
```

```
#
```

```
#stat file.txt
```

```
File: file.txt
```

```
Size: 0          Blocks: 0          IO Block: 4096   regular empty file
```

```
Device: 10303h/66307d  Inode: 16961354   Links: 1
```

```
Access: (0664/-rw-rw-r--)  Uid: ( 1000/ec2-user)   Gid: ( 1000/ec2-user)
```

```
Access: 2024-07-13 05:22:35.308253378 +0530
```

```
Modify: 2024-07-13 05:22:35.308253378 +0530
```

```
Change: 2024-07-13 05:22:35.308253378 +0530
```

```
Birth: 2024-07-13 05:22:35.308253378 +0530
```

```
#
```

Which Means..

More Writes..

More Writes..

Per Unit of Data..

The Challenge Of Reliable Storage

- Disks have **Disk Controllers** and they have **Caches INSIDE OF THEM**
- Disk IOPS Are Slow
- Most databases rely on filesystems and not use disks directly
- And filesystems have their own metadata to maintain
- These filesystems are mounted in operating systems
- And operating systems have **THEIR OWN Page Cache**

----- Databases

----- Operating System (with it's cache)

----- Filesystem

----- Disks (with it's Caches)
(outside the control of OS)

Cache Writing Policies

- Write-Back
- Write-Through

Linux WriteBack Mechanism

```
# ps -ef | grep writeback  
root      25      2  0 Jun28 ?          00:00:00 [writeback]
```

Linux Background Kernel Thread

How does a database guarantee a write ?

Linux Write-Through Mechanism

`fsync()` and `fdatasync()`

Linux Syscalls

How does it help in guaranteeing a write ?

`fsync()` and `fdatasync()`

CRUD

C - Create

R - Read

U - Update

D - Delete

Databases need to handle a lot of CUD..

And each CUD can result in multiple writes..

How ?

If `fsync()` and `fdatasync()` are synchronous operations, and disk writes take long

Then how do databases perform fast with the high CUD operations

Without compromising on efficiency and reliability.. ?

Write Ahead Log

Different storage solutions refer WAL by different names

- Write Ahead Log (WAL) in Postgresql in Ceph SDS, RabbitMQ
- Binlog in MySQL
- Transaction Log in ElasticSearch, MongoDB, Neo4j
- Journals in most filesystems (xfs, ext4)
- Append Only File (AOF) in Redis
- Commit Log in Cassandra
- Redo Log in Oracle

Problems solved for CUD

`fsync()` and `fdatasync()` solved reliability

WAL solved efficiency in writes

But what about efficiently organising writes ?

Let's talk about storing the data on disk for
efficient reads..

The common ways data is organised

LSM

B-Tree

A word about indexing

Back To ElasticSearch..

Storage Layers in Elasticsearch

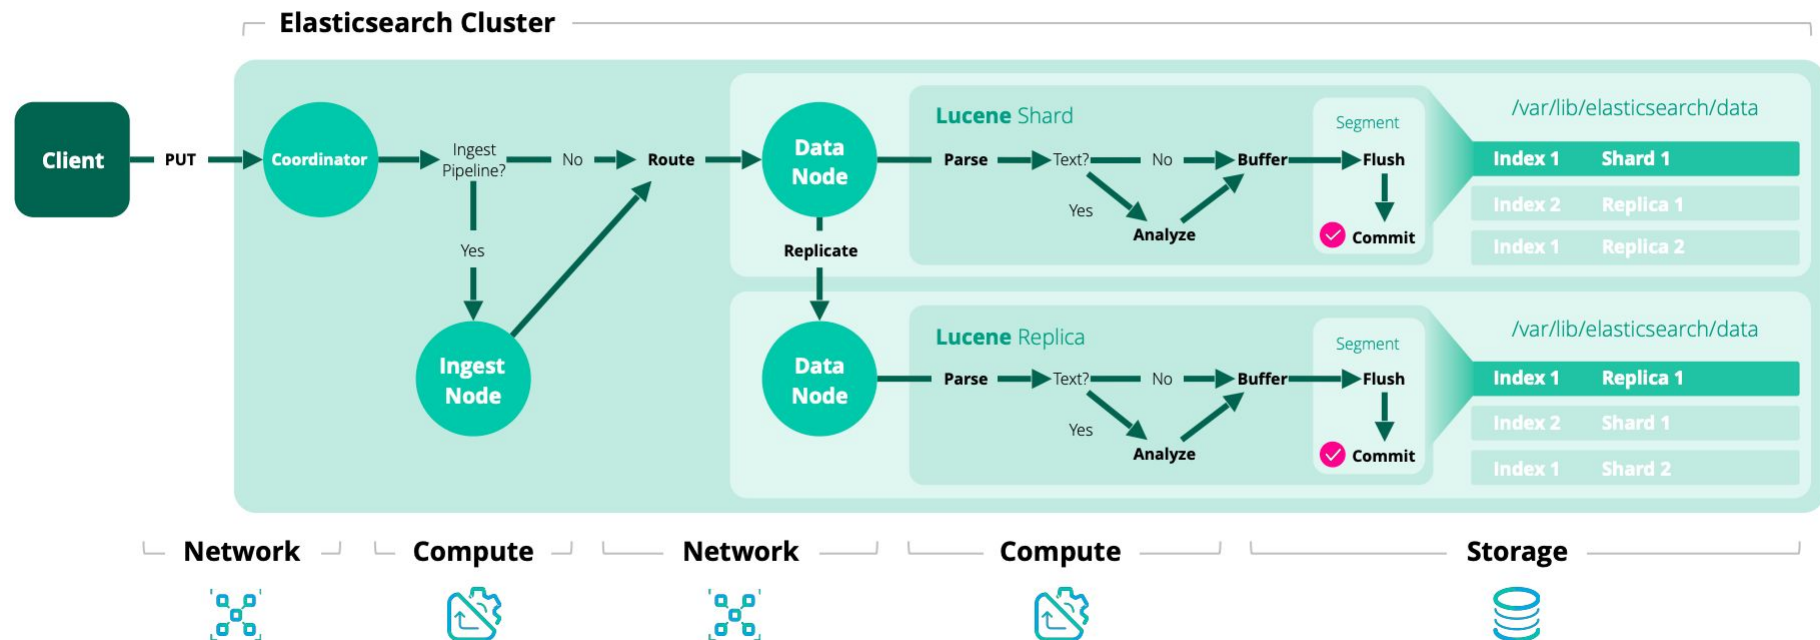
- Segment Files
- Shards
- Index
- Alias
- Data Streams

Components in the Elasticsearch storage

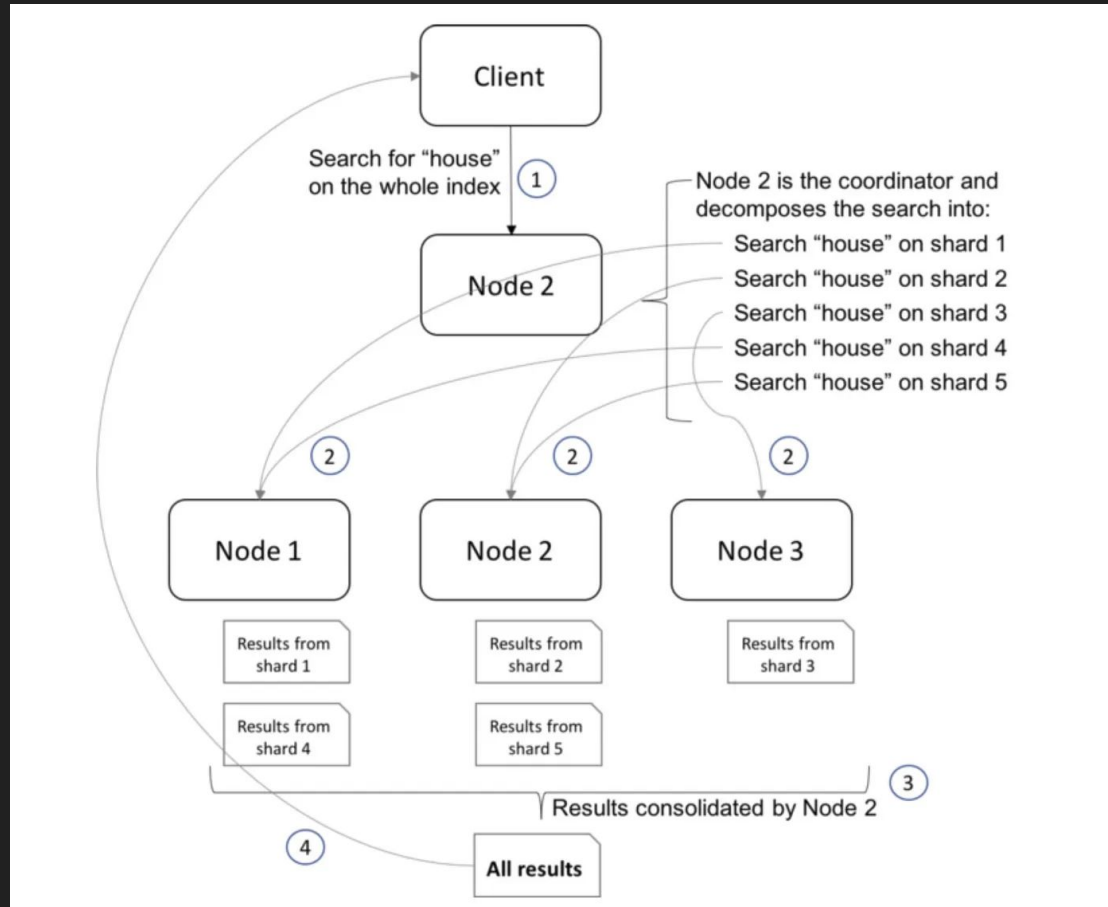
- In-Memory Buffer
- Segment Files
 1. Committed Segment Files
 2. Searchable Segment Files
- Segment Files Merge
- Tombstones

Data Flow Overview

Write Path Overview



Read Path Overview



Write Path Encounters




Refresh

- Can be thought of as a lightweight fsync
- Happens every second by default in Elasticsearch
- Flushes the in-memory buffers to create an in-memory segment and allow it to be searched

Flush

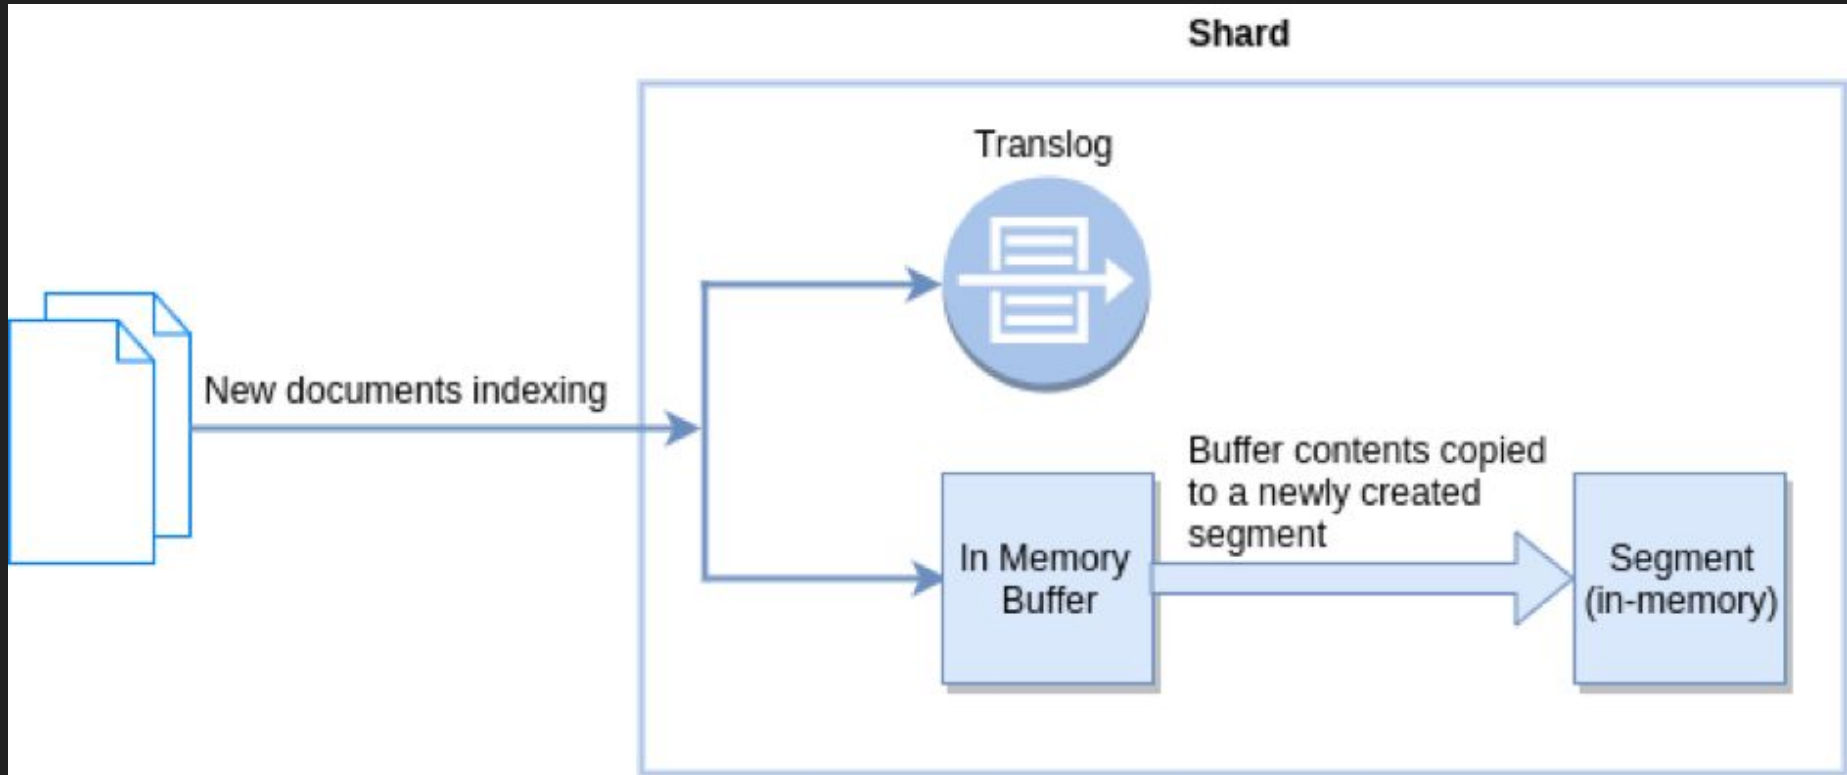
- An actual fsync to merge the segments and flush the merged segments on to disk

Refresh And Flush



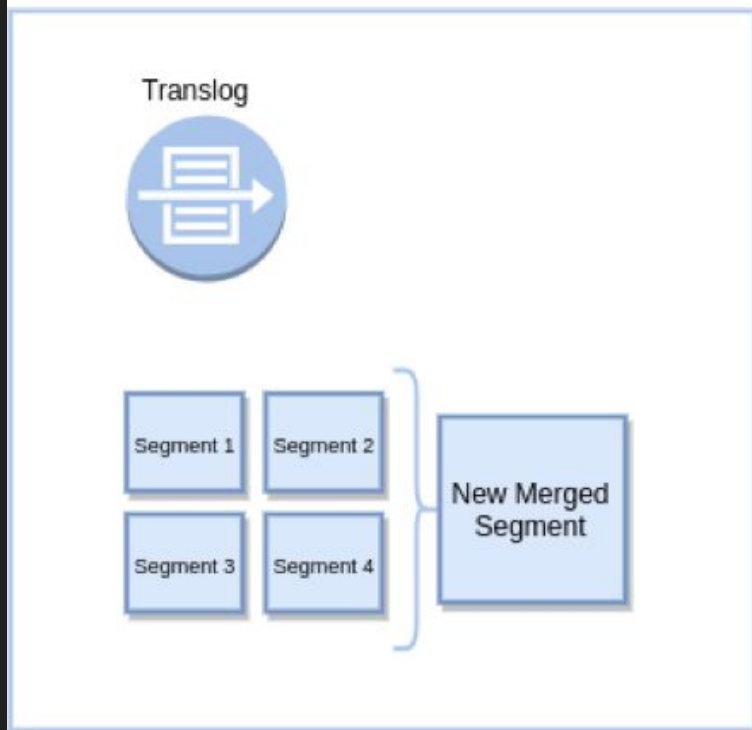
Yet Another Diagrams to Explain Refresh And Flush

What Refresh Does



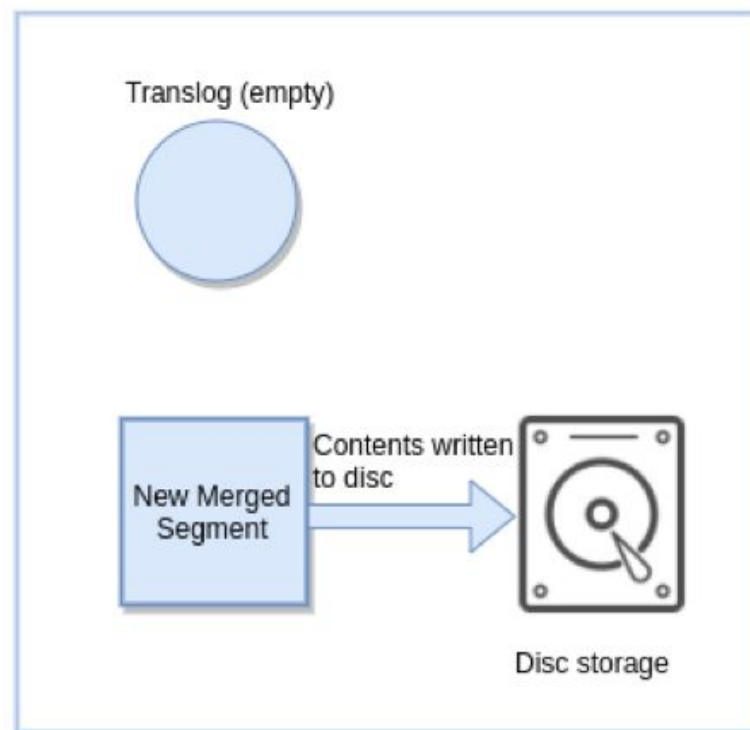
What Flush Does

Shard



Stage 1
Figure 3

Shard



Stage 2
Figure 4

Index and Indexing are not the same..

Understand the context..

Types of Indexing a field can have

- Text (for full-text search)
- Keyword (for exact match)

Understanding how the R in the CRUD Works in Elasticsearch

The power of vocalisation

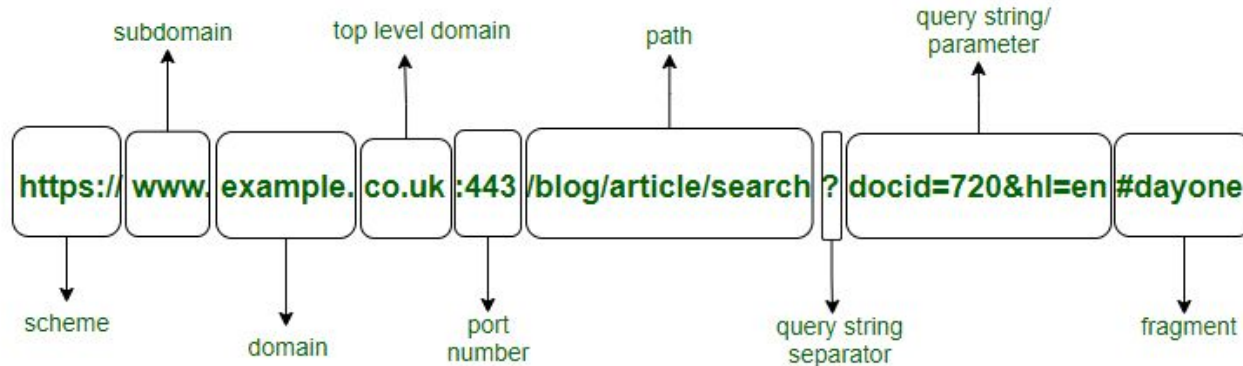


Gossip Protocol

Structure of a URL

Parts of a URL

URL : <https://www.example.co.uk:443/blog/article/search?docid=720&hl=en#dayone>



Examples of a few HTTP APIs

GET /_cat/indices

```
# curl https://192.168.1.10:9200/_cat/indices?v
```

health	status	index	uuid	pri	rep	docs.count	docs.deleted	store.size	pri.store.size
green	open	biz.live.data.2024.04.04-000024	fCmpWN2ySt6zZisx9Qj0vA	2	1	7	0	258.5kb	258.5kb
green	open	biz.live.data.2024.05.04-000025	Zo27FF_VT5-Wqi0mBa-DsA	2	1	9	0	293.7kb	293.7kb
green	open	biz.live.data.2024.06.03-000026	IIq6DbwZTbKdRg3g9IqMcw	2	1	44	0	582.8kb	582.8kb
green	open	biz.live.data.2024.07.03-000027	WX3GgeLUQ_qwzPzCZQEgRg	2	1	19	0	949.8kb	474.9kb

GET /_cat/aliases

```
# curl https://192.168.1.10:9200/_cat/aliases?v&s=alias,index
```

alias	index	filter	routing.index	routing.search	is_write_index
biz.live.data	biz.live.data.2024.04.04-000024	-	-	-	false
biz.live.data	biz.live.data.2024.05.04-000025	-	-	-	false
biz.live.data	biz.live.data.2024.06.03-000026	-	-	-	false
biz.live.data	biz.live.data.2024.07.03-000027	-	-	-	true

GET /_cat/nodes

```
# curl https://192.168.1.10:9200/_cat/nodes?v
```

ip	heap.percent	ram.percent	cpu	load_1m	load_5m	load_15m	node.role	master	name
192.168.1.10	57	93	12	1.18	1.12	1.01	mdi	*	node-1
192.168.1.11	53	94	7	0.47	0.69	0.77	mdi	-	node-2
192.168.1.12	54	91	5	0.16	0.27	0.35	mdi	-	node-3
192.168.1.13	55	90	7	0.34	0.21	0.27	di	-	node-4
192.168.1.14	52	90	6	0.74	0.74	0.42	di	-	node-5

GET /_cat/allocation

```
# curl http://192.168.1.10:9200/_cat/allocation?v
```

shards	disk.indices	disk.used	disk.avail	disk.total	disk.percent	host	ip	node
1063	118gb	120.2gb	179.7gb	1tb	83	192.168.1.10	192.168.1.10	node-1
1061	132.9gb	136.3gb	163.5gb	1tb	85	192.168.1.11	192.168.1.11	node-2
1061	153.6gb	157.4gb	142.4gb	1tb	87	192.168.1.12	192.168.1.12	node-3
1063	122.6gb	127.1gb	172.7gb	1tb	84	192.168.1.13	192.168.1.13	node-4
1063	199.8gb	103.2gb	196.6gb	1tb	82	192.168.1.14	192.168.1.14	node-5

GET /_cluster/health

```
# curl http://192.168.1.10:9200/_cluster/health?pretty
{
  "cluster_name" : "es-cluster-1",
  "status" : "green",
  "timed_out" : false,
  "number_of_nodes" : 5,
  "number_of_data_nodes" : 5,
  "active_primary_shards" : 5310,
  "active_shards" : 10622,
  "relocating_shards" : 0,
  "initializing_shards" : 0,
  "unassigned_shards" : 0,
  "delayed_unassigned_shards" : 0,
  "number_of_pending_tasks" : 0,
  "number_of_in_flight_fetch" : 0,
  "task_max_waiting_in_queue_millis" : 0,
  "active_shards_percent_as_number" : 100.0
}
```

Other helpful APIs in Elasticsearch

- GET `/_cat/recovery?v&active_only=true`
- GET `/_cat/thread_pools?v`
- GET `/_cluster/settings?flat_settings=true&include_defaults=true`
- GET `/_cat/shards?v`
- GET `/_cat/segments?v`

A word about managing data at large scale

In Elasticsearch

Tiered Storage, Searchable Snapshots
Index Priorities, ILM, SLM policies

Signoff

@GypsyCosmonaut

<https://www.linkedin.com/in/gypsyscosmonaut/>