

OpenTelemetry - The Universal Observability Umbrella



About me



Abhinav Kapoor

- Principal Engineer, Nagarro
- Work as a Technical Architect for .NET Solutions.
- 15+ years of technical experience in architecting & developing solutions for Banking, Supply chain & Embedded domains.
- I enjoy writing about software architecture, communication patterns, modernization & more on Medium.com <https://medium.com/@kapoorabhinav>.



Accreditations



I'm available on Linked at
<https://www.linkedin.com/in/abhinav-kapoor-0394456/>

What we'll cover



1. Observability Concepts

1. What is observability & Why do we need It ?
2. How does it fit in a cloud native landscape?
3. Application Instrumentation – The pillars for getting insights
 1. Logs
 2. Metrics
 3. Traces

2. Future of Observability - OpenTelemetry

1. What is it & why is it needed?
2. Its components.
 1. Traces
 2. Metrics
 3. Logs

3. Demo

What is observability ?



Sir Lewis Carl Davidson Hamilton
- 7-time F1 champion



Steering Wheel with controls and dashboard
- Image credit Mercedes X

What is observability ?



- Moving from Formula 1 to Modern Software applications
 - Complex systems – Lot of moving parts
 - High stakes
 - Cutting edge engineering
 - Fast Reaction
 - Several teams with different specializations at work
- Observability is the ability to observe what is going on inside the system.
- It can be considered as an NFR which helps measures other NFRs

Why do we need Observability?



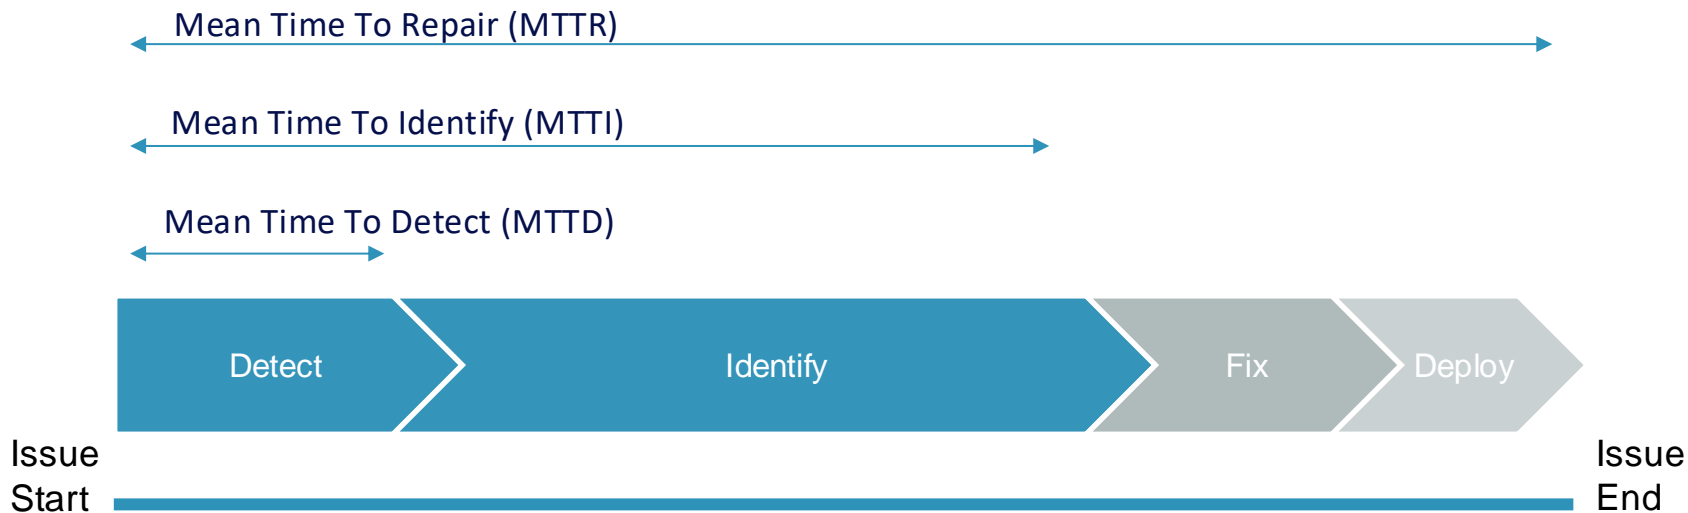
- We want to avoid these reactions



Why do we need Observability?




- Issue Timeline in Operation Support



Why is Observability more critical now?



- Evolution to smaller but more services - high cohesion and low coupling.




“Complexities arises when the dependencies among the elements become important”
- We have reduced code complexity but introduced much more moving parts


Why is Observability more critical now?



Technology transformation trends



Microservices



Polyglot Applications




Ephemeral & cloud native environments

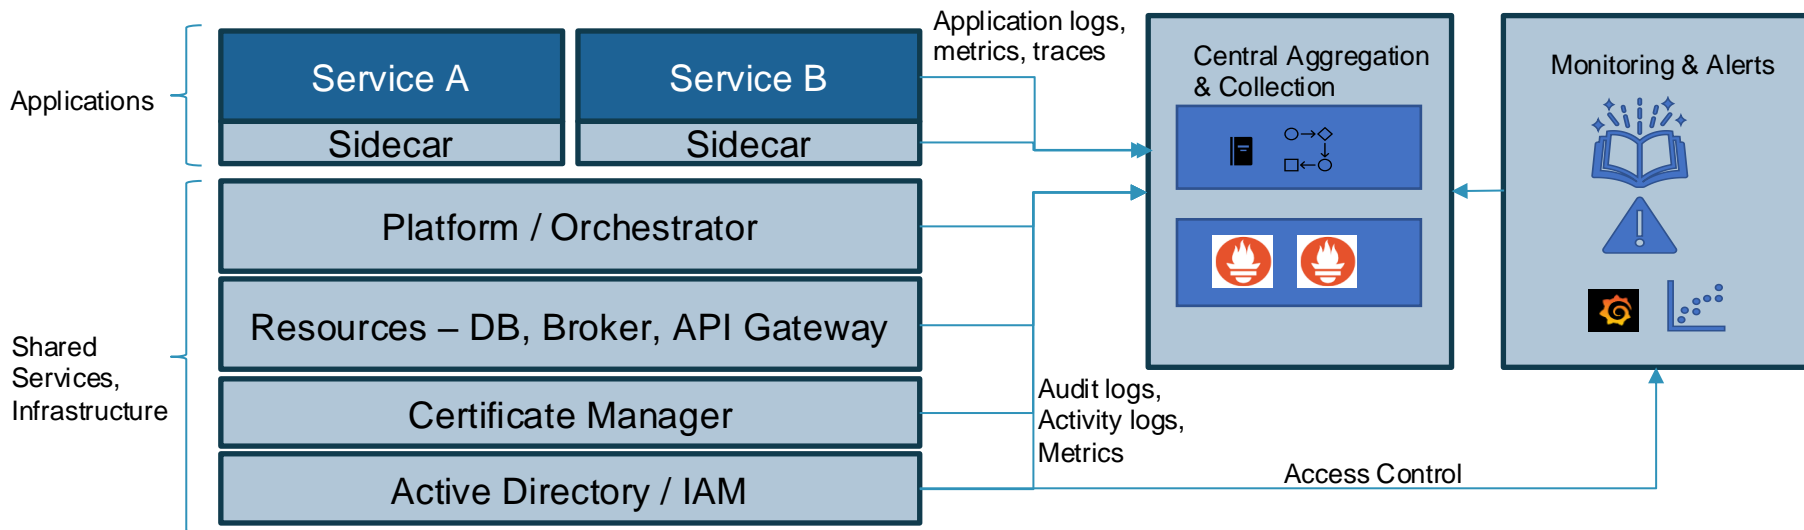


How does it fit in a cloud native landscape?

How does it fit in a cloud native landscape?



How does it fit in a cloud native landscape?





Deep Dive - Application Instrumentation

Application Instrumentation




- Logs – What happened & who did it ?
 - Important events. State changes, who did.
- Traces – How did it happen ?
 - Flows within the program or outside the program.
- Metrics – How much happened ?
 - Numbers which can be aggregated to have formulas – SLOs, KPI & alerts
- Health – How is the system doing ? Current state of health

Big picture



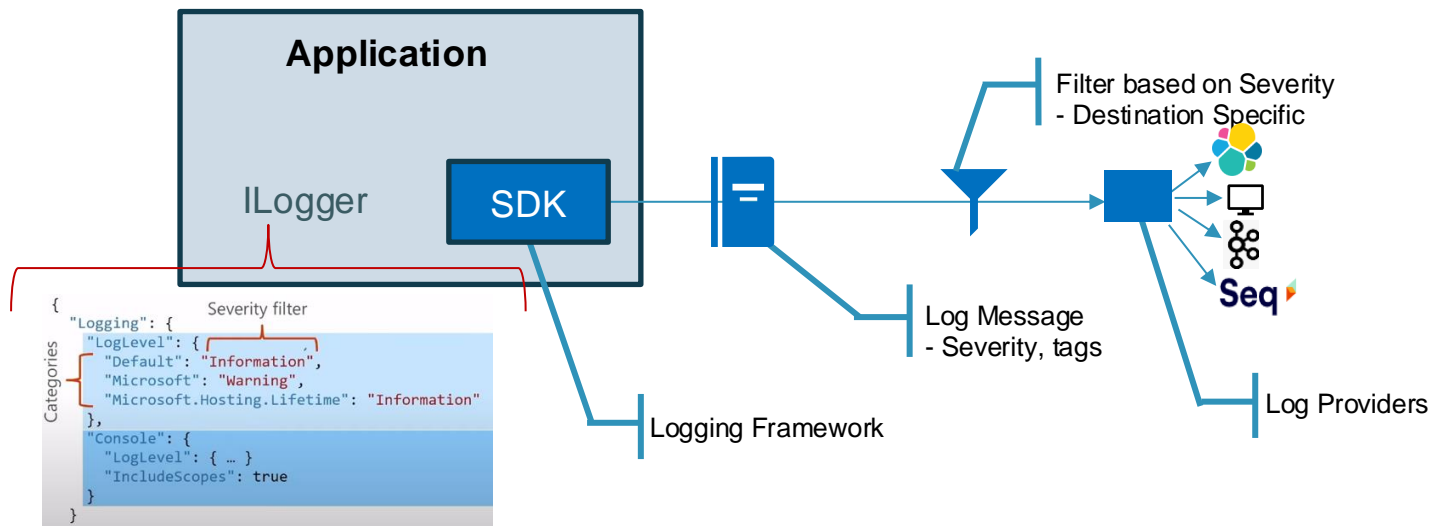
- We don't just want to have connected black boxes.
- While infrastructure & sidecars can give a good overview, there are always blind spots.
- Do an **intentional** instrumentation.
- It speeds up development.





Logs – What Happened ?

Logs – Components of Logging

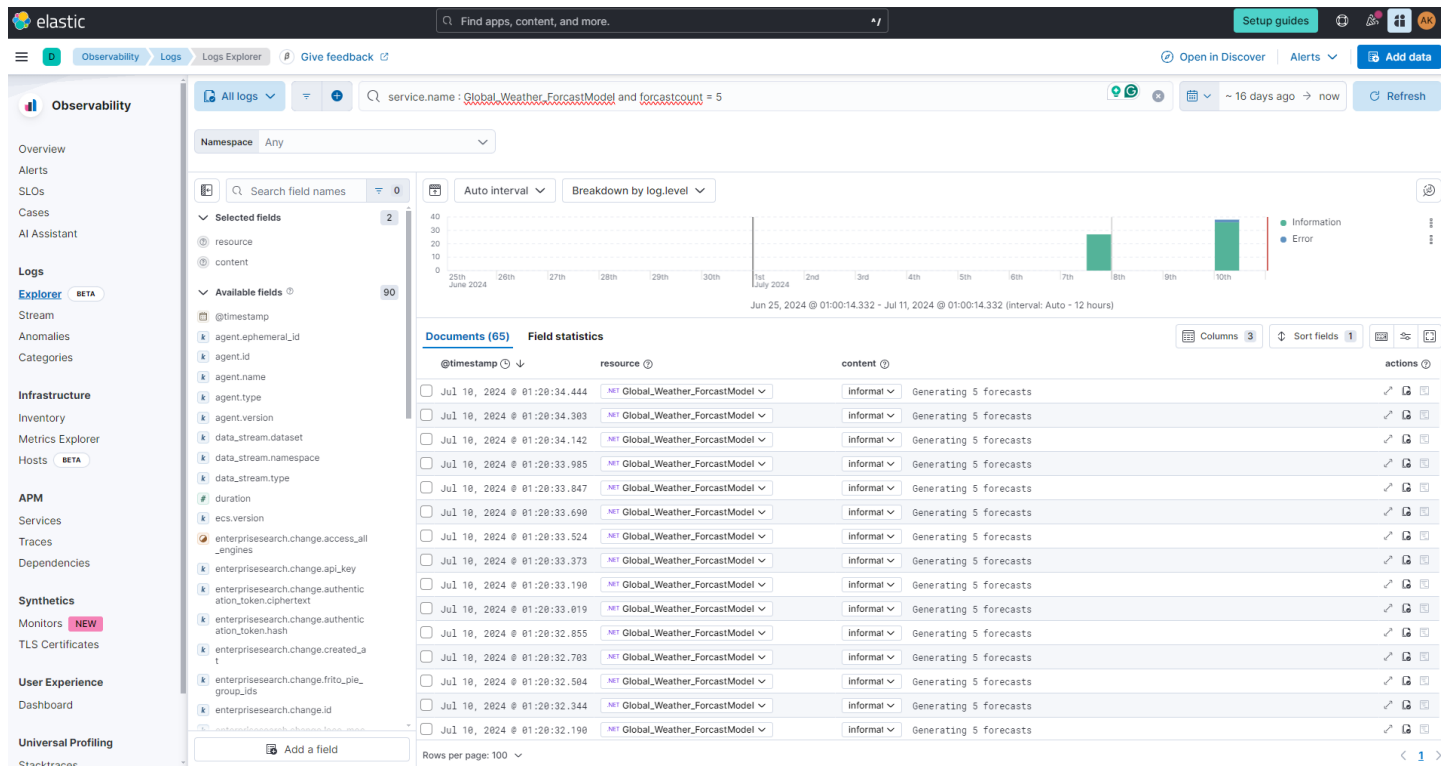


Types of Logs



- Based on Information they contain – Governs where they are stored, how long they are stored, who can access them
 - Technical – Startup configurations, system event that the application receives
 - Business - Business and User Information
- Based on the schema
 - Unstructured – Simple statements with severity.
 - Structured / Schematic logs - Logs with Metadata to provide context. Pod, Service, IP address & more.

Duality of Logs - Readable messages & Quarriable metrics





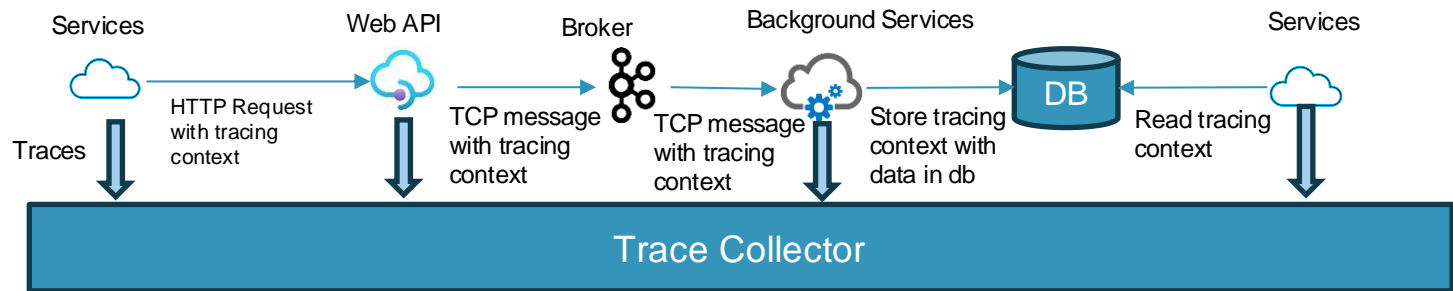
Traces - How did it happen ?

Traces – Types of traces



- Tracing – This is similar to logging. But noisier and collects more information from deeper parts of the application.
- Distributed Tracing – It's the method to observe requests as they propagate through different services.
 - Its structured log which comes from different processes, nodes and can be stitched together to give an end to end view.

Distributed Tracing

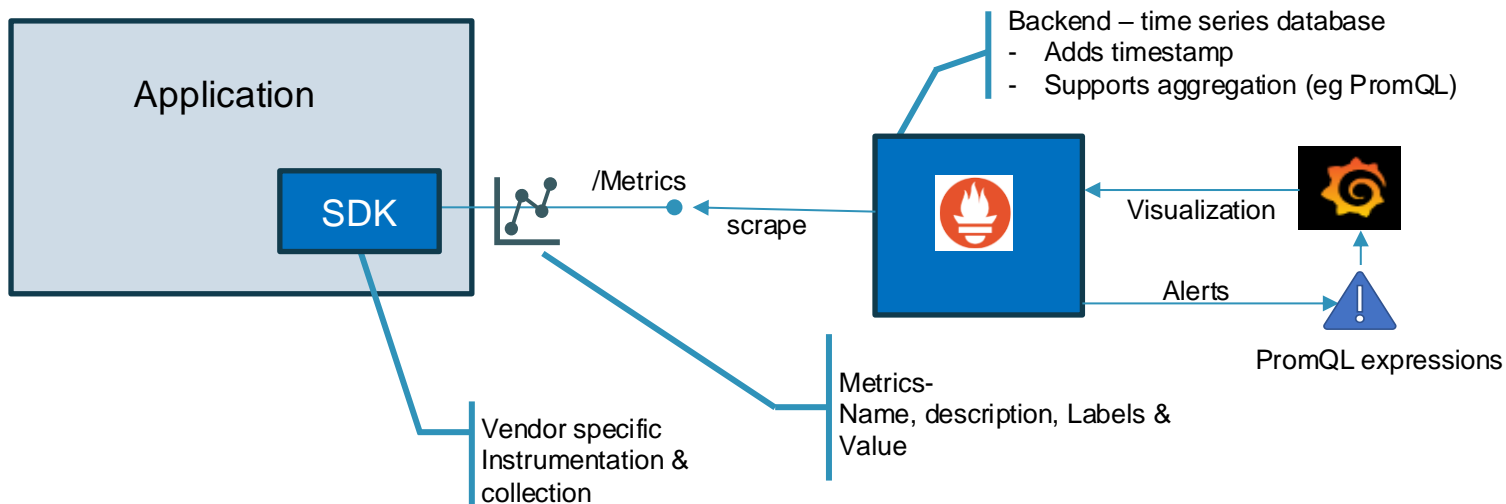


- Correlation IDs are passed in the remote calls/messages.
- The receiver regenerates context proceeds with its unit of work.
- Depending upon business transaction, traces can be of few seconds to hours or days.



Metrics—How much is happening?

Metrics – Components



Types of Prometheus Metrics - Counters



Track the occurrence of an event in application.

- The absolute value may not be helpful.
- The delta between two timestamps or the rate of change over time is helpful. Functions like `Increase()`, `Rate()`.
- Consider the case when applications restart.
- **Examples - Orders processed counter, unsuccessful counters, business error, technical errors.**
- Sample output of meta-endpoint (`/metrics`)

```
# HELP http_requests_total Total number of http api requests
# TYPE http_requests_total
counter http_requests_total{api="add_Product", company="ITC"} 3433
counter http_requests_total{api="add_Product", company="TATA"} 2000
```

PromQL

```
increase(http_requests_total{api="add_product", company="ITC"}[5m])
```

Types of Prometheus Metrics - Gauges



Snapshots of a metric at a single point in time. Not an event.

- **Example is message queue size, CPU utilization at a given time.**
- Useful functions - `avg_over_time`, `max_over_time`, `min_over_time`, and `quantile_over_time`

Types of Prometheus Metrics - Histograms



- Histograms divide the entire range of measurements into a set of intervals—named buckets—and count how many measurements fall into each bucket.
- These buckets are defined at compile time, they are upper inclusive (le).
- Histogram looks like the following

```
http_request_duration_seconds_sum{api="add_product" instance=" host1 "} 8953.332
http_request_duration_seconds_count{api="add_product" instance=" host1"} 27892
http_request_duration_seconds_bucket{api="add_product", instance=" host1", le="0.05"} 1672
http_request_duration_seconds_bucket{api="add_product", instance=" host1", le="0.1"} 8954
http_request_duration_seconds_bucket{api="add_product", instance=" host1", le="0.25"} 14251

sum by (le) (rate(http_request_duration_seconds_bucket[5m]))
```

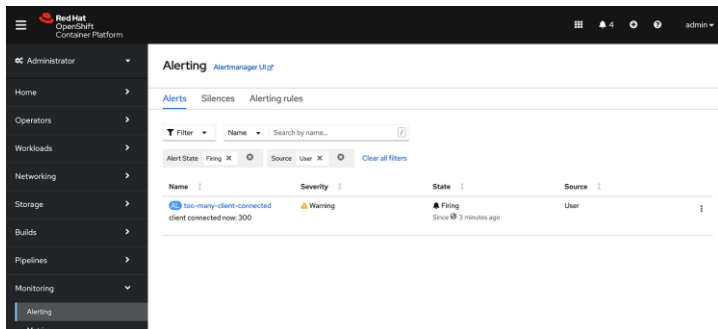
Types of Prometheus Metrics - Summaries



- Similar to Histogram but they add quantile label to bound the summary.
- OpenTelemetry has been marked it legacy.
- Summary looks like the following
 - `http_request_duration_seconds_sum{api="add_product" instance="host1.domain.com"} 8953.332`
 - `http_request_duration_seconds_count{api="add_product" instance="host1.domain.com"} 27892`
 - `http_request_duration_seconds{api="add_product" instance=" host1 " quantile="0"}`
 - `http_request_duration_seconds{api="add_product" instance="host1" quantile="0.5"} 0.232227334`
- Considerations due to quantile.

Metrics - What happens after instrumentation?

- The orchestrator can enrich the application instrumentation with infrastructure information such as Service Name, Pod, IP address.
- The backend saves it with timestamp
- Add Alerts & dashboards.

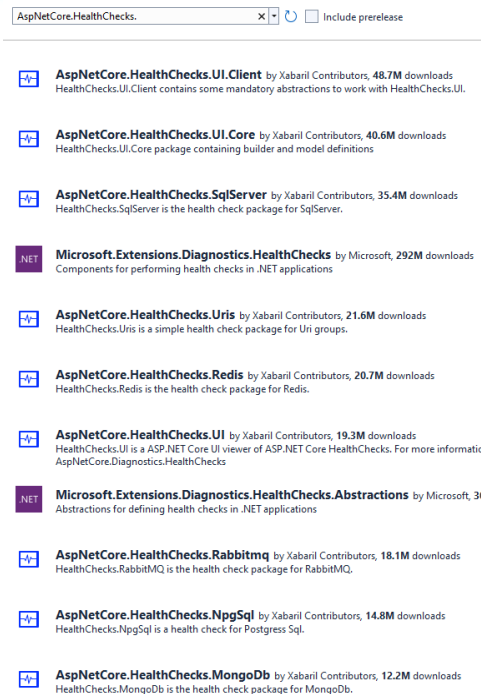




Health Check– How is the system doing ?

Health check

- Basic use is liveness probe.
- More useful when we add dependencies.
- Prebuilt libraries depending upon the language
- Considerations
 - Orchestrator may restart the service while the issue is in a dependency.
 - Checking all downstream services could be time consuming.
 - Meta-end points (/health) may not be exposed. Instead of writing to the response metrics may be used to indicate issue.





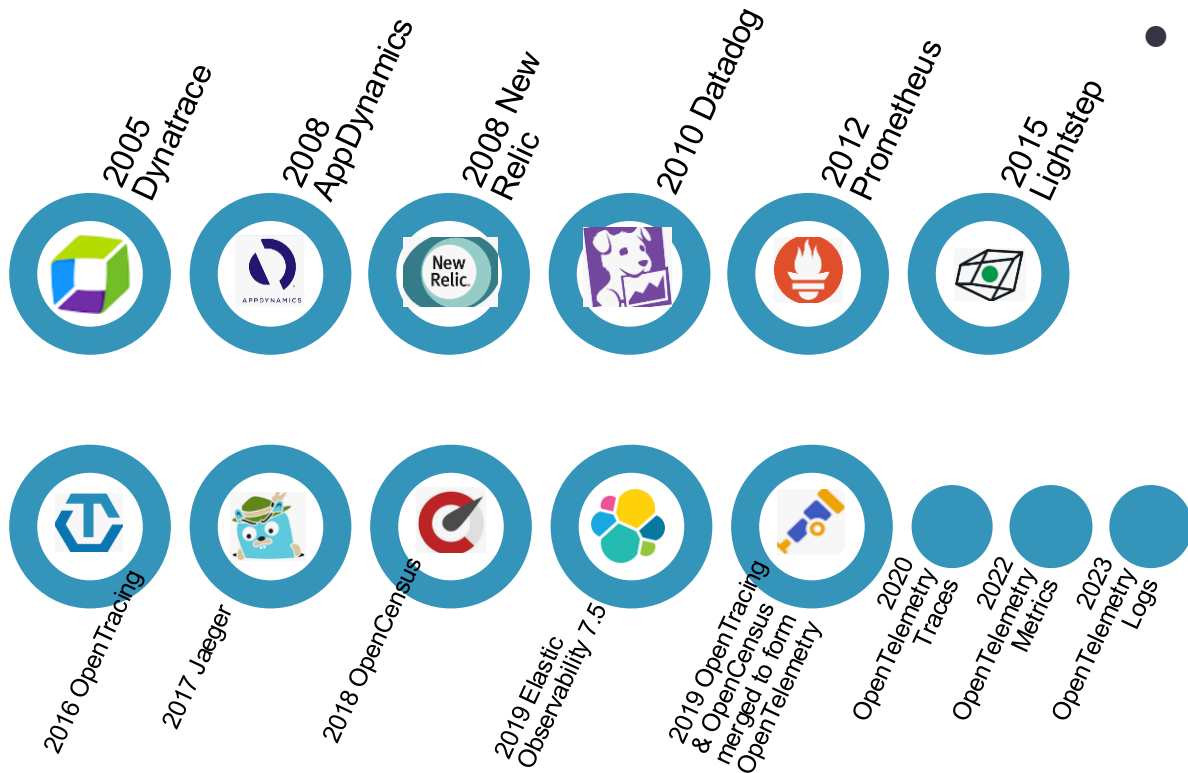
OpenTelemetry

What is OpenTelemetry ?



- It is an open-source observability framework for **infrastructure & application instrumentation**.
- It's a Protocol at heart which aims to be **vendor agnostic**. It produces SDK in multiple languages.
- As Open Telemetry doesn't provide a backend implementation (its concern is **creating, collecting, and sending signals**), the data will flow to another system or systems for storage and querying.
- 3 Signals -
 - Traces - 2020 – Capture distributed traces & propagate context.
 - Metrics 2022 – Capture Metrics from application & infrastructure.
 - Logs 2023 – Performant Logs with unified semantics.

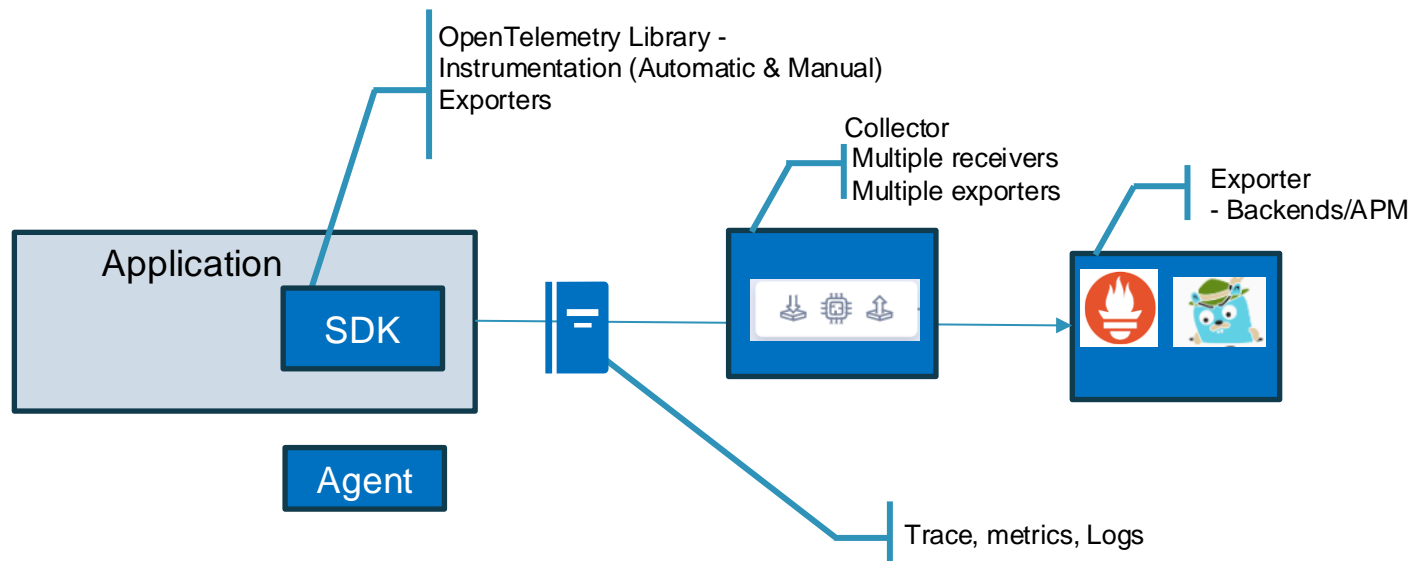
Why OpenTelemetry?



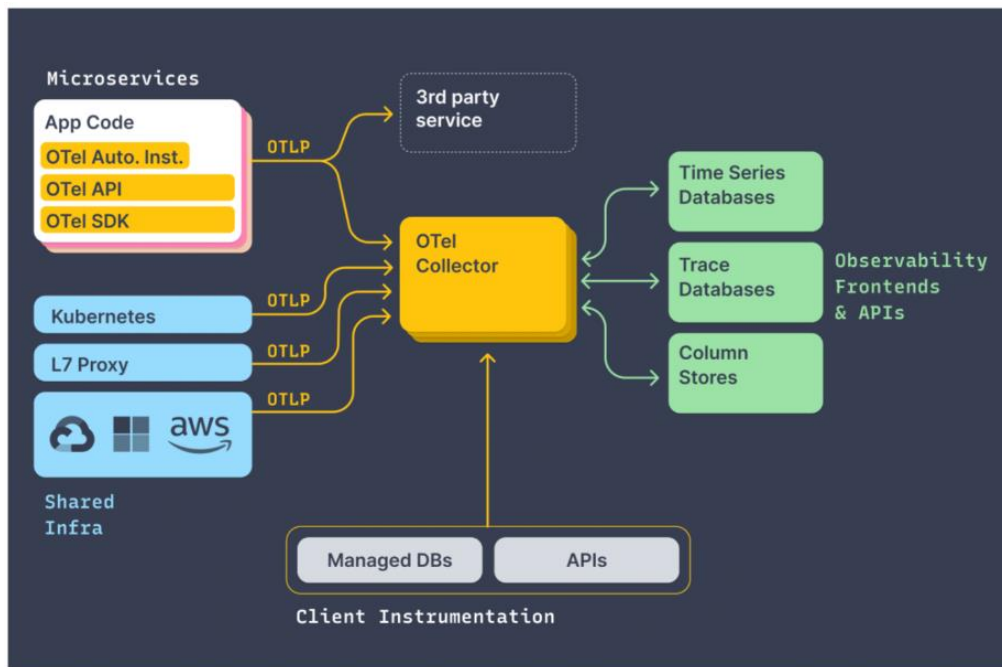
● Before Open Telemetry

- APM vendors had their own things – SDKs, Agents, Grammar
- Vendor lock-in.
- Application developers, built wrappers, event hooks to keep application isolated.
- Competing Open standards – OpenTracing (Traces), OpenCensus (Metrics, Traces)
- Different Standards for Logs

Components of OpenTelemetry



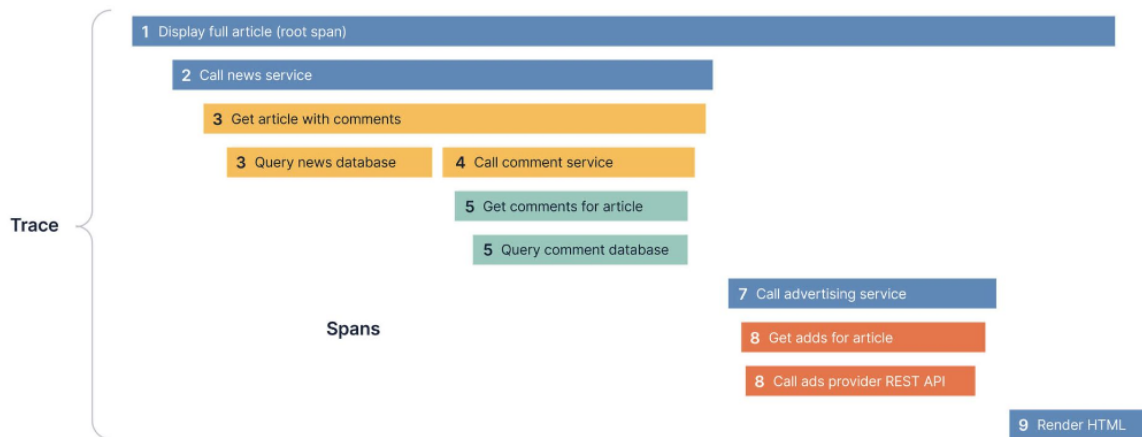
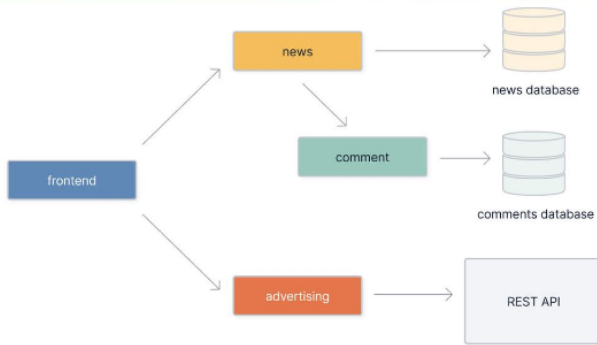
Components of OpenTelemetry





OpenTelemetry – Traces & Spans

Traces & Spans



- Trace is the “Thing” being done – Usecase
- The “Thing” is composed of smaller things either locally or remotely (called Spans)
- It’s a container for Spans

Traces & Spans




- Span is time bound analysis
- Scoped piece of work
- It is a structured blob of data
 - Uniqueld (SpanID)
 - CorrelationId (TraceId)
 - CausalityID (parentSpanID)
 - Point in time action (Events)
 - Related Spans (Links – Causal vs Casual)
 - Enrich with Tags Attributes



Traces & Spans



- Spans can be in process or out of process
- Keep Spans wide than deep



Signals of OpenTelemetry –Logs & Metrics



Structured logs

(All) Level: (All) Filters: TraceId == 1f283e509c9c4bfcca17724802e45daf

Resource	Level	Timestamp	Message	Trace	Details
328b16cc-2541-4e90-90e...	Information	10:54:21.867 AM	Data Saved	1f283e5	View
328b16cc-2541-4e90-90e...	Information	10:54:21.878 AM	Forecast updated to Balmy temperature -14 for 07/12/2024 10:54:21	1f283e5	View

Total: 4 results found

Log entry details GlobalWetherStation.SatelliteConsumer

Resource: Global_Weather_ForecastModel Timestamp: 10:54:21.878 AM

Log entry 5 ^

Name	Value
Level	Information
Message	Forecast updated to Balmy temperature -14 for 07/12/2024 10:54:21
datetime	12-07-2024 10:54:21
temperature	-14
summary	Balmy

Context 3 ^

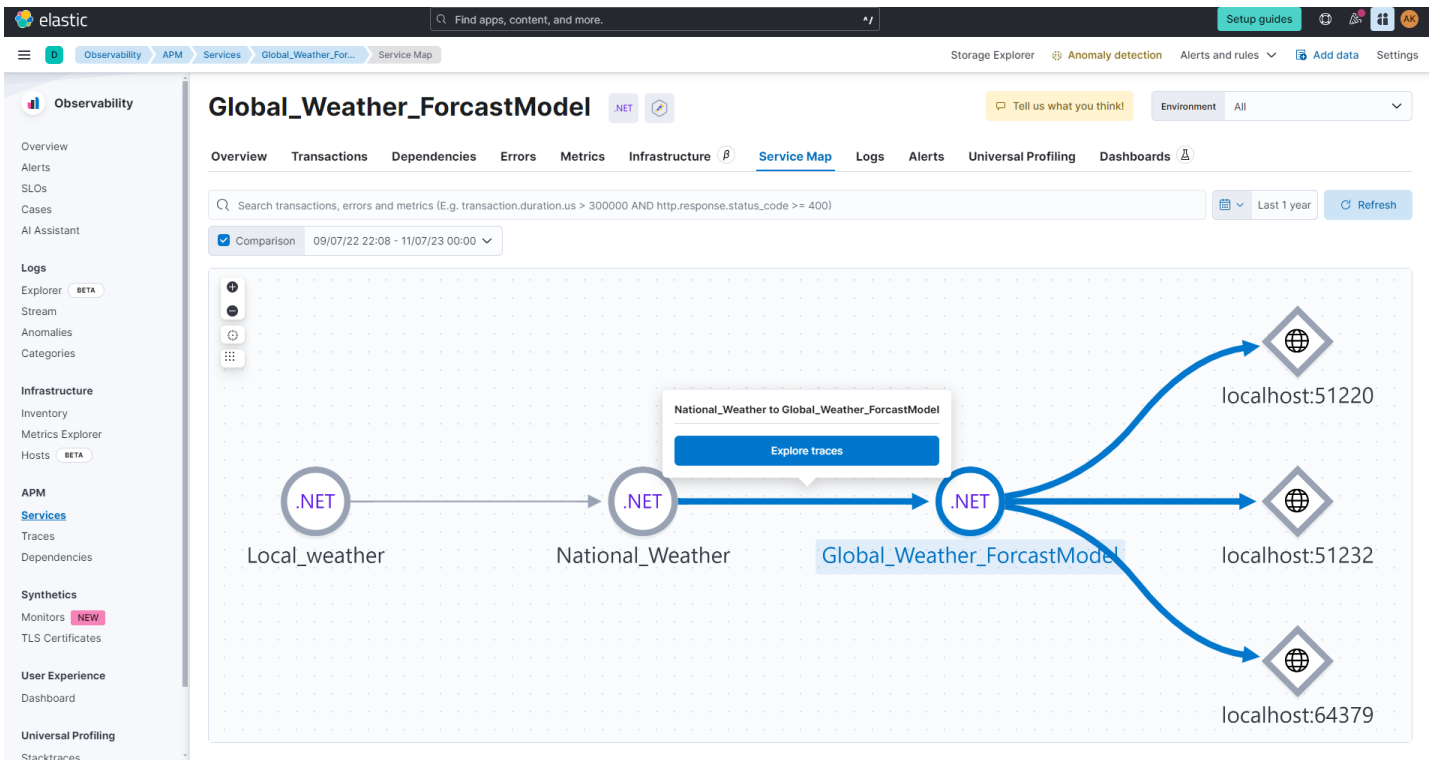
Name	Value
Category	GlobalWetherStation.SatelliteConsumer
TraceId	1f283e509c9c4bfcca17724802e45daf
SpanId	238e22bd667f439f

- Related to Spans
- Logs may be kept for several years due to legal reasons (unlike traces)
- Summary is marked deprecated



Demo

Demo





Questions ?



Thank you!