

### Introduction



Hi, I am **Abhishek Garg**. Currently working as *DoE* in **Syfe**.

Today I will be talking about one of architecture project we did back in 2018, when prometheus was getting pace and Thanos was very new.

I will talking about how we did migrate our observability from a enterprise to a in-house monitoring system which eventually helps us save millions of USD / Year.

#### Let's talk about Scale first !!

Nodes / Endpoints	Data Ingestion sum(scrape_samples_scraped)	Retention	Cost (DataDog)
25k	300 Million / minute	30-60 Days	\$ 6 Million / year *

#### \*Datadog Current Pricing

- APM → \$31 \$45 (range reflects different tiers and access to features) / host / month
- Infrastructure Monitoring → \$15 \$34 (range reflects different tiers and access to features) /host /month

Current Base price around **USD 14 million / Year** 



# What is Prometheus??

### What it is for

#### **Metrics-based monitoring & alerting stack.**

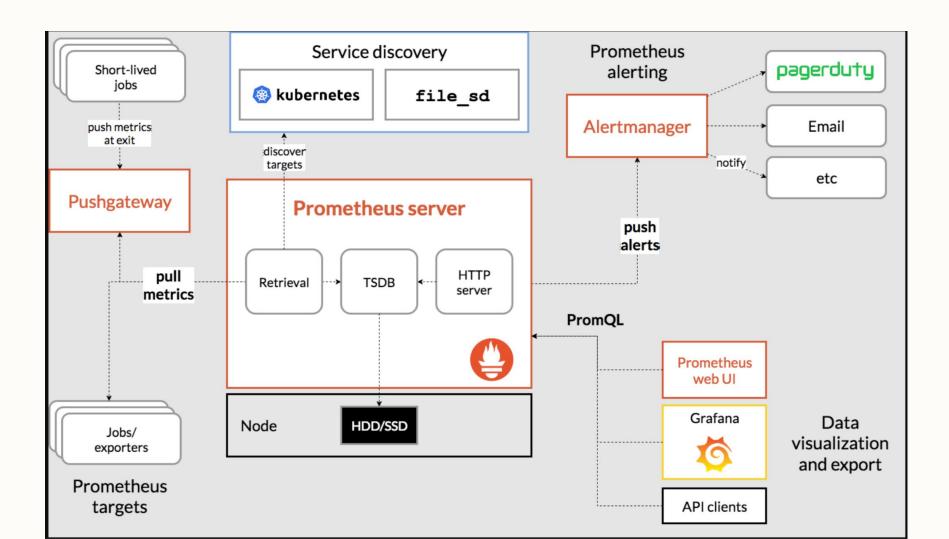
- Instrumentation
- Metrics collection and storage
- Querying, alerting, dashboarding
- For all levels of the stack!

### What it is not for

#### Logging, Tracing and ....

- Logging, Event Collection or tracing
- Automatic anomaly detection
- Scalable or durable storage
- Automatic Horizontal scaling
- User authorization management





It's Simple, right?



### Not Much, here are some problems!

#### 1. Scalability Issue

- a. Single Server Limitations: Prometheus is designed as a single-node system, which can become a bottleneck when dealing with high cardinality metrics and large-scale environments.
- b. Storage Limitations: The local storage of Prometheus can be a limiting factor for long-term storage and high-frequency data points.

#### 2. High Cardinality and Performance

- b. High Cardinality Metrics: Prometheus can struggle with high cardinality metrics (metrics with a large number of unique label combinations), which can lead to increased memory usage and degraded performance.
- c. Query Performance: Complex queries over large datasets can be slow and resource-intensive, impacting the responsiveness of the system.

#### more...

#### 3. Configuration Management

a. Static Configuration: Prometheus relies on static configuration files for setting up scrape targets and rules, which can be cumbersome to manage in dynamic and large-scale environments.

#### 4. Operational Overhead

b. Maintenance and Upgrades: Maintaining Prometheus, including handling updates, scaling, and troubleshooting, can require significant operational effort and expertise.

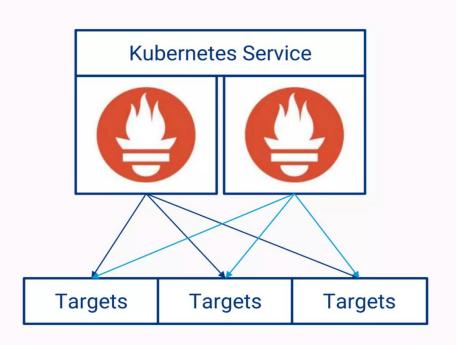
Let's talk business!



### **Problems**

**Solutions** 

### Highly Available Prometheus!!!

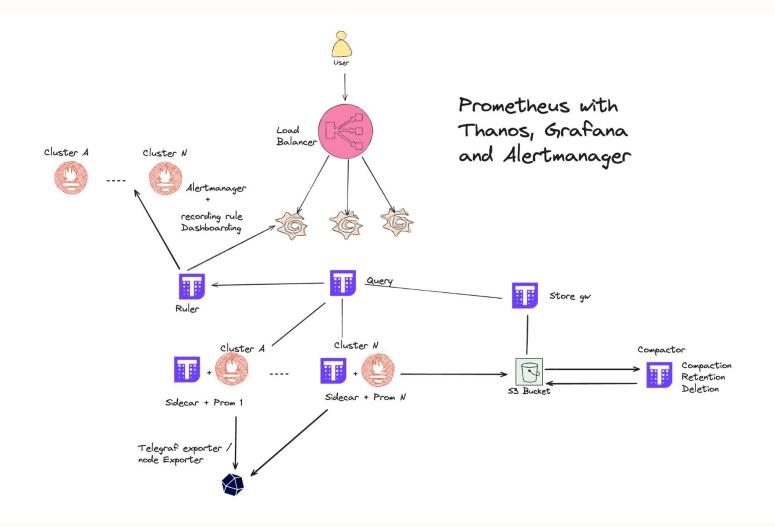


#### Not without its challenges:

 When you refresh the data, you will see it change as metrics will potentially differ between the two instances



Thanos is a set of components that can be composed into a highly available metric system with unlimited storage capacity, which can be added seamlessly on top of existing Prometheus deployments.



### Seems Costly?

·
.
.

Let's Calculate!

Components Used	Thanos, Prometheus, Grafana, Object Storage, ELB	
RAM requirement for <b>300 million / min</b> metrics	Approx 7000 GB	
Storage requirements for above estimates metrics	46 TB ( 46000 GB)	
Cost of Graviton Instance (1 Year Upfront) / GB RAM / year	36.2 USD	
S3 Cost / GB	0.08 (assuming 10000 RW call / GB)	

Compute total Cost = ( 7000 \* 36.2 ) = 253400 USD

Storage total Cost = (46000 \* 0.08 \* 12 ) = 44160 USD

Final Cost = Compute + Storage = ( 253400 + 44160 ) = 2,97,560 USD / Year

#### Savings:

#### **Savings Percentage:**

$$Savings~Percentage = \left(\frac{Savings}{Datadog~Cost}\right) \times 100 = \left(\frac{5702440}{6000000}\right) \times 100 \approx 95.04\%$$

#### **Summary**:

Savings Percentage: Approximately 95.04%

**Note:** even with a **45**% variance added as noise to the cost of Prometheus + Thanos, you would still save approximately **92.81**% compared to using Datadog.

#### Other Important points:

- Avoid Unnecessary Labels and metrics: Minimize the use of labels that can have a large number of unique values, such as timestamps, user IDs, session IDs, or request IDs. If possible try to combine them. Also drop excessive metrics or limit metric collection / host.
- **Limit Label Values:** Ensure that labels have a limited and predictable set of values. For example, instead of using exact URLs as labels, use URL patterns or endpoints.
- Aggregate Metrics: Use Prometheus recording rules to pre-aggregate metrics with high cardinality.
   This reduces the amount of data stored and queried. For example, aggregate metrics by time intervals or other meaningful dimensions.
- **Downsampling:** Apply downsampling techniques to reduce the resolution of historical data, retaining only the necessary level of detail.
- **Optimized PromQL Queries:** Write efficient PromQL queries that avoid scanning large datasets unnecessarily. Use functions like sum, avg, max, min, and rate to aggregate data effectively.
- **Query Caching:** Implement query caching mechanisms if possible to reduce the load on Prometheus when executing repeated queries.
- Monitoring and Alerting: Set up alerts to notify you when cardinality exceeds acceptable thresholds, enabling proactive management.

### Meet the Team



Ram Shankar Jaiswal
Sr. Cloud Architect



Abhishek Garg
Director of Engineering

# Q & A

## **Thanks**

Scan here to connect me over LinkedIn

