

Observability – Cost Vs Value

OpenTelemetry signals optimization

Traditional Operations - Old School

Operations, as a key entity within businesses, has typically functioned independently, often gaining prominence after deployment and billed according to resource usage

- Operations are generally seen as separate revenue model.



Cloud Era - Pay as you use model

- **SaaSification** has transformed the way applications are delivered, shifting from on-premise to the cloud and as a service model.
- With a "**Pay As You Use**" charging system, dev, devOps, and operations are integrated, making Operations equally accountable for the company's bottom line.
- Optimization and applying **FinOps** to reduce operations costs are now **shared responsibilities**.
- **Everyone takes ownership for their cloud usage**



Designed by [Freepik](#)

FinOps - Basic principles

- Use the “**Pay As You Use**” model to your advantage
- Know the business value
- Accessibility of Data
- Optimize
- Collaborate, collaborate & collaborate



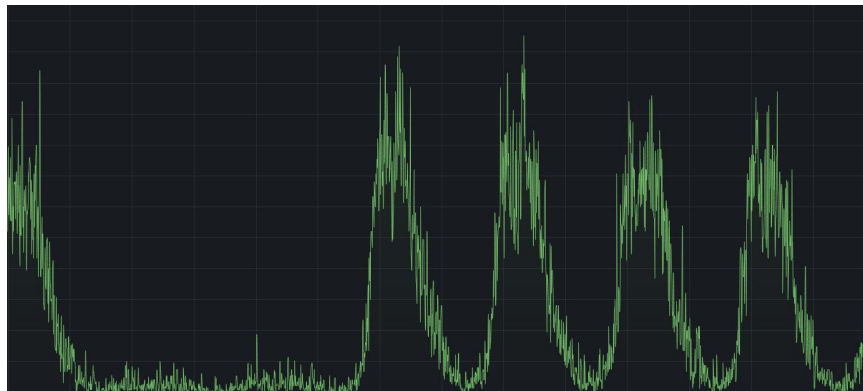
Designed by [Freepik](#)

Observability with OpenTelemetry and Grafana

Traces - Distributed trace gives you a hierarchy for calls within an API. Individual call within the trace is called **span**.

Metrics - Metric is a time series representation of data points related to an attribute. Example average response time of an API over last 2 days

Logs - Logs are the streams of output from the application.

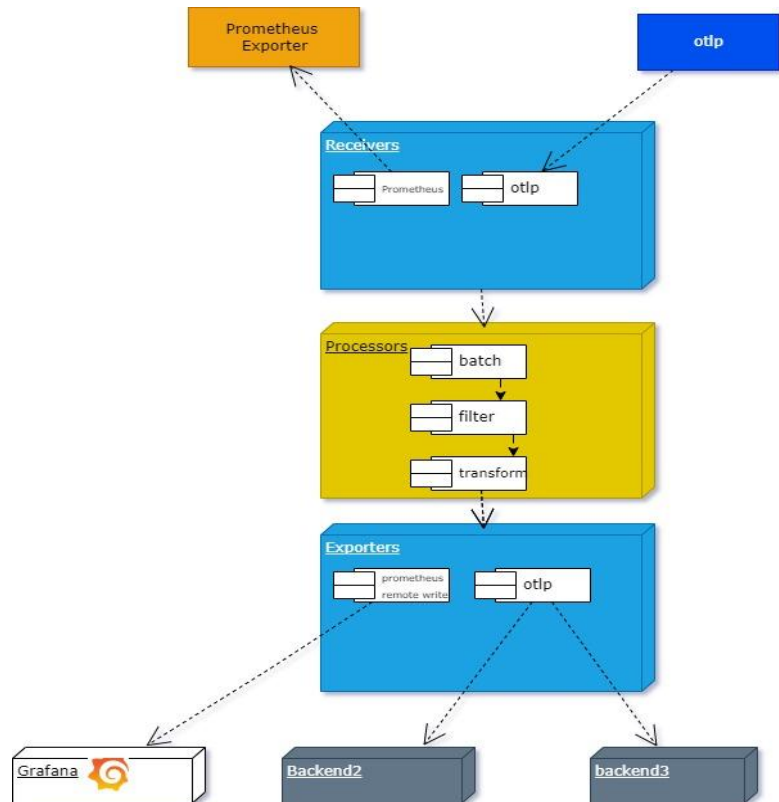


OTEL Collector workflow

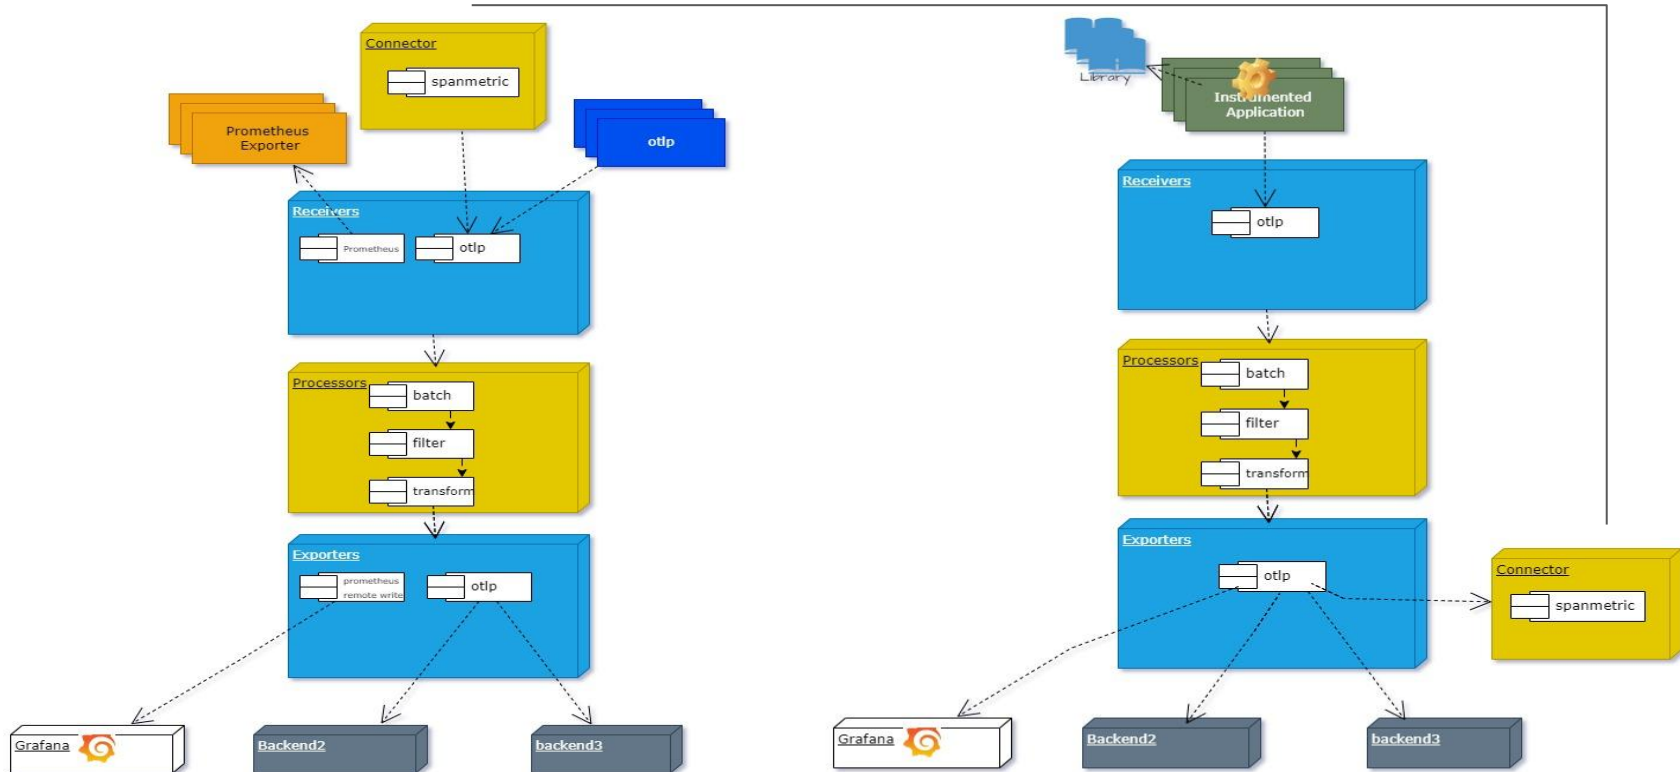
Receivers : definition of agents/process which will receive/pull the signal data.

Processors : Define the configuration to enrich, filter or transform signal data.

Exporters : definition of process to push the signal data to a storage medium. Backend storage can accept OTLP schema or exporter can convert it to their native schema.



Trace and Metric workflows



Observability - Signals and their use model

Traces

- Based on the **size**
- Base on the **Memory size** monitored
- Based on the **count**

Metrics

- Based on the **number of series**
- Based on the **number of attributes** being monitored
- Based on **size**

Logs

- Based on **size**

Observability - Know the business value

Traces

- Need for **sampling**
- Understanding of **Key APIs**
- Understanding of various services and components
- Call Patterns

Metrics


- Need for retention.
- Histogram **Granularity**
- Aggregation
- Samples per minute
- Required **labels**

Logs

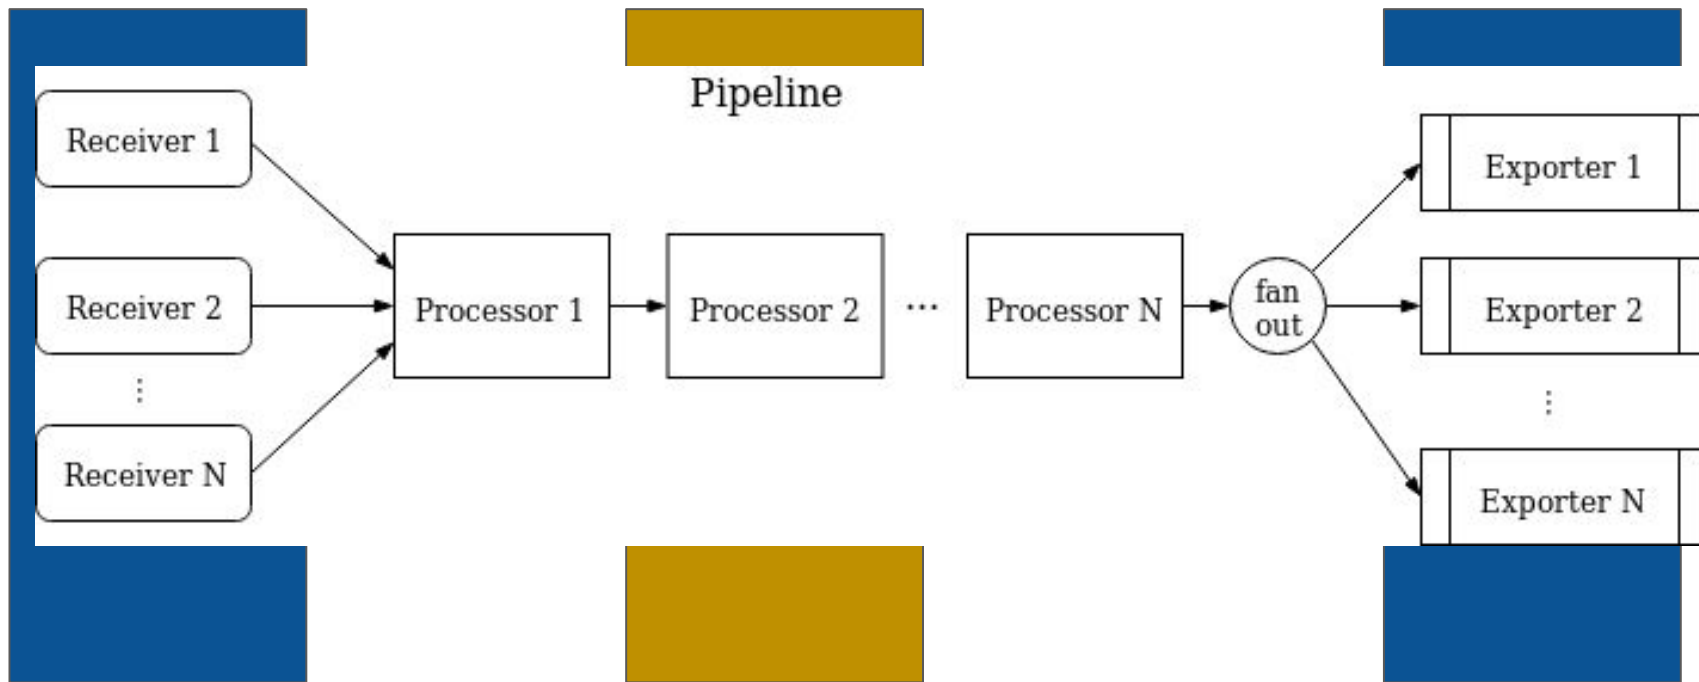
- Structure and **log pattern**
- Log Levels
- **Filtering** needs

Observability - Accessibility of usage metrics

- Based on the billing criteria, usage data need to be accessible in **real time**.
- Allow you to **drill down** the usage on various parameters.
- Alerts can be setup for any **anomaly** or usage beyond a limit



Cost Optimization in open Telemetry



Tips - Sampling at receiver level - (Trace)

Sampling at the source level allow you to process only selected traces based on specific conditions.

- **Head Sampling** - Sample a percentage of traces.
- **Tail Sampling** - Sample based on the condition and criteria within the trace

```
export OTEL_TRACES_SAMPLER="traceidratio"  
export OTEL_TRACES_SAMPLER_ARG="0.1"
```

```
import io.opentelemetry.sdk.trace.SdkTracerProvider;  
import io.opentelemetry.sdk.trace.samplers.Sampler;
```

```
public class Example {
```

```
    public static void main(String[] args) {
```

```
        // Configure the tracer provider with the desired sampler
```

```
        SdkTracerProvider tracerProvider = SdkTracerProvider.builder()
```

```
            .setSampler(Sampler.alwaysOn()) // Set to always sample traces // or
```

```
            .setSampler(Sampler.alwaysOff()) // Set to never sample traces // or
```

```
            .setSampler(Sampler.traceIdRatioBased(0.5)) // Set to sample a fraction of tr
```

```
            .build();
```

```
    }
```

```
}
```

Tips - Enabling and Disabling libraries

Many a time traces get **duplicated** as there are multiple layers of instrumentation available. In such cases one layer can be disabled.

- Enabled by default, review the traces and disable one at a time.

```
OTEL_INSTRUMENTATION_[NAME]_ENABLED=false
```

```
OTEL_INSTRUMENTATION_COMMON_DEFAULT_ENABLED = false
```

Tips - Filter Processor (trace, metric, logs)

Filter processor filter the incoming signal based on the otel query(**open telemetry transformation language** aka otel) and discard the selected spans before sending to exporters.

processors:

filter/otl:

error_mode: ignore

traces:

span:

- 'attributes["container.name"] == "app_container_1"'
- 'name == "app_3"'

spanevent:

- 'attributes["grpc"] == true'
- 'IsMatch(name, ".*grpc.*")'

metrics:

metric:

- 'name == "my.metric" and resource.attributes["my_label"] == "abc"'

datapoint:

- 'metric.type == METRIC_DATA_TYPE_SUMMARY'
- 'resource.attributes["service.name"] == "my_service_name"'

logs:

log_record:

- 'IsMatch(body, ".*password.*")'



Tips - Transform Processor (trace, metric, logs)

Transform processor allow you to apply various functions based on the query filtering through otel. It allows applying the functions like to **rename, replace, keep, set, truncate** etc.

transform:

error_mode: ignore

trace_statements:

- context: resource

statements:

- keep_keys(attributes, ["service.name", "service.namespace", "cloud.region", "process.command_line"])

- replace_pattern(attributes["process.command_line"], "password\\=[^\\s](\\s?)", "password=***")*

- limit(attributes, 100, [])

- truncate_all(attributes, 4096)

- context: span

statements:

- set(status.code, 1) where attributes["http.path"] == "/health"

- set(name, attributes["http.route"])

- replace_match(attributes["http.target"], "/user//list/*", "/user/{userId}/list/{listId}")*

- limit(attributes, 100, [])

- truncate_all(attributes, 4096)

Tips - Transform Processor - Continue

transform:

error_mode: ignore

metric_statements:

- context: resource

statements:

- keep_keys(attributes, ["host.name"])

- truncate_all(attributes, 4096)

- context: metric

statements:

- set(description, "Sum") where type == "Sum"

- convert_sum_to_gauge() where name == "system.processes.count"

- convert_gauge_to_sum("cumulative", false) where name ==

"prometheus_metric"

- context: datapoint

statements:

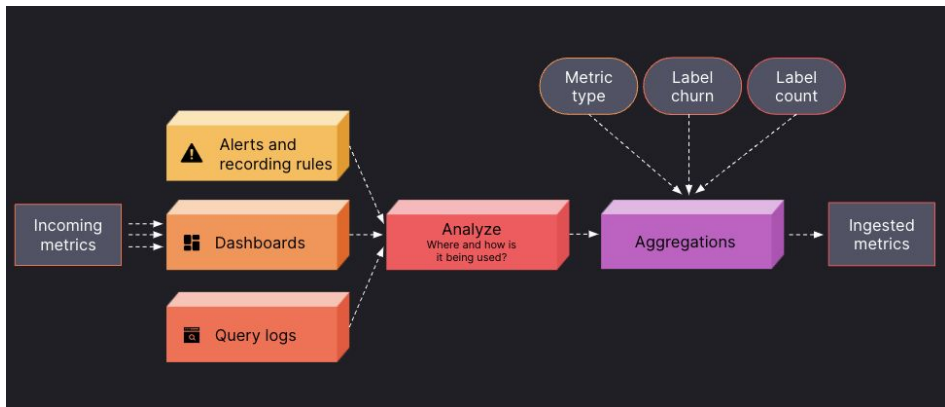
- limit(attributes, 100, ["host.name"])

- truncate_all(attributes, 4096)

Tips - Aggregation at exporter/backend level

Adaptive Metrics (Grafana)

- Analyse all the use of any metric (dashboard, alerts, queries) and suggest metrics which are **unused** for filtering.
- Suggest aggregation to **reduce data points**





Designed by [Freepik](#)

Tips and Tricks - Other Ways

- Drop the Labels
- Reduce data point
- Expire Serieses

Observability - Know the business value

- Traces

Samples
or all the
calls ?

All API
or Key
ones ?

Limit the
size of the
value ?

All
services
or
selected
one ?

Club calls
based on
pattern?

Observability - Know the business value

- Metrics

Retention
Period ?

Series
aggregation?

Histogram
Granularity ?


Any
duplication
?

What all
labels are
needed?

Observability - Know the business value

- Logs

-



Quiz Question

Q1: Everyone knows K8s is the biggest CNCF project but Which project of CNCF has the 2nd highest velocity ?

A: OpenTelemetry

Q2: What are the 3 major signals of Observability ?

A : Traces, Logs and Metrics

Q3 : What are the major section/clauses of Otel Collector ?

A: Receiver, processor & exporter

Q4 List 3 use model for Observability tools ?

A, Size, Number of series, RAM size

Q5: Name processors which help in finOps for Observability ?

A: Filter, transform

